



DeepLearning.AI

Data Storage and Queries

Storage Systems

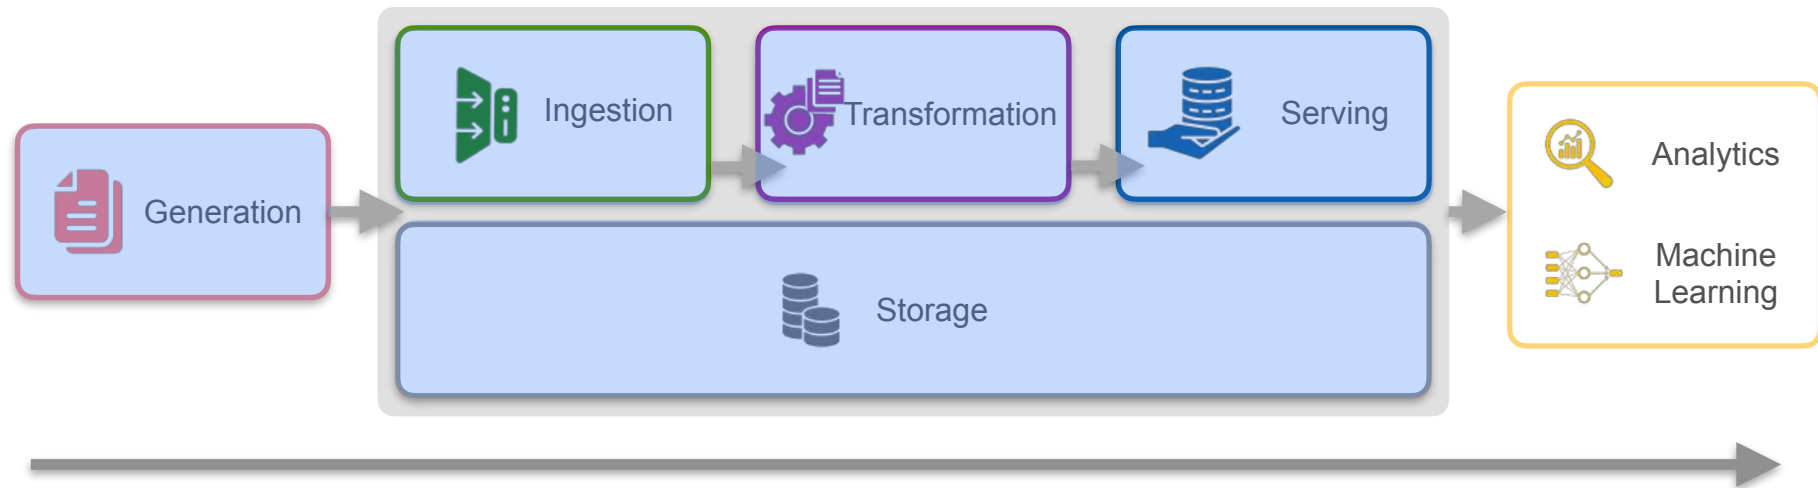


DeepLearning.AI

Storage Systems

Course 3 Overview

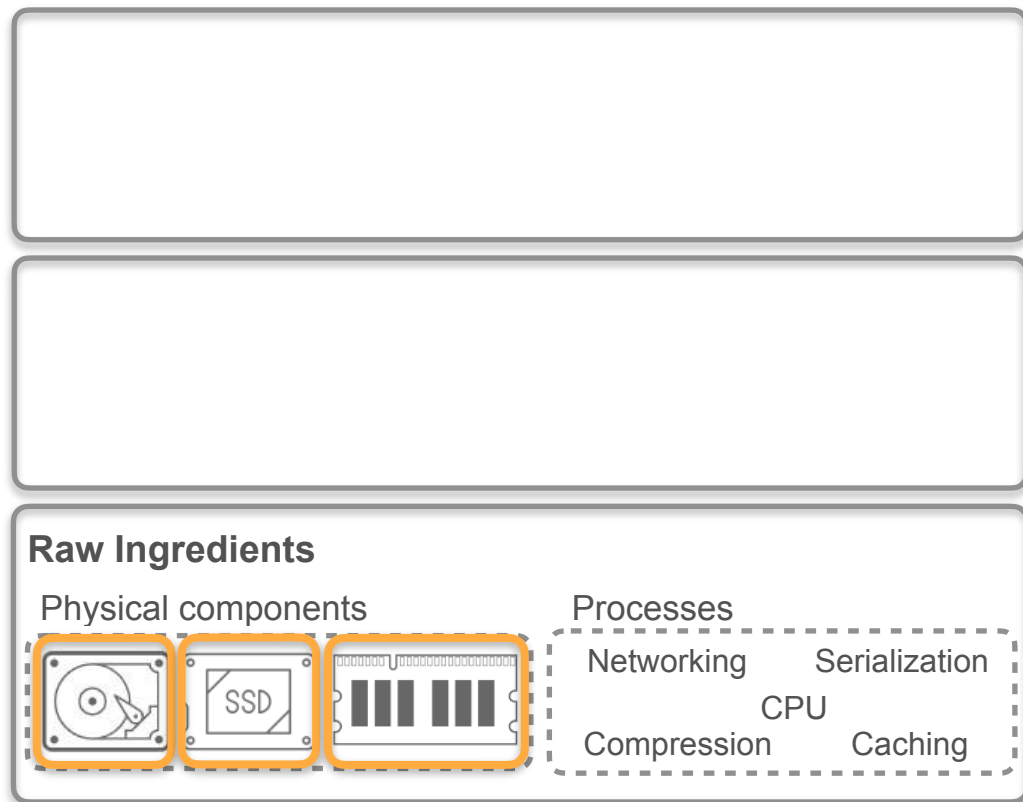
Storage



Storage solution considerations:

- Data type
- Data size
- Data format
- Access and update pattern

Storage Hierarchy



Storage Hierarchy

Management system:

Organizes data in the raw components and allows you to interact with the stored data



OLTP Systems

Online Transactional Processing Systems
Focus on performing read and write queries with low latency

OLAP Systems


Online Analytical Processing Systems
Focus on applying analytical activities on data (e.g. aggregation, summarization)

Storage Systems



Raw Ingredients

Physical components



Processes

Networking

Serialization

CPU

Compression

Caching

Storage Hierarchy

Management system:

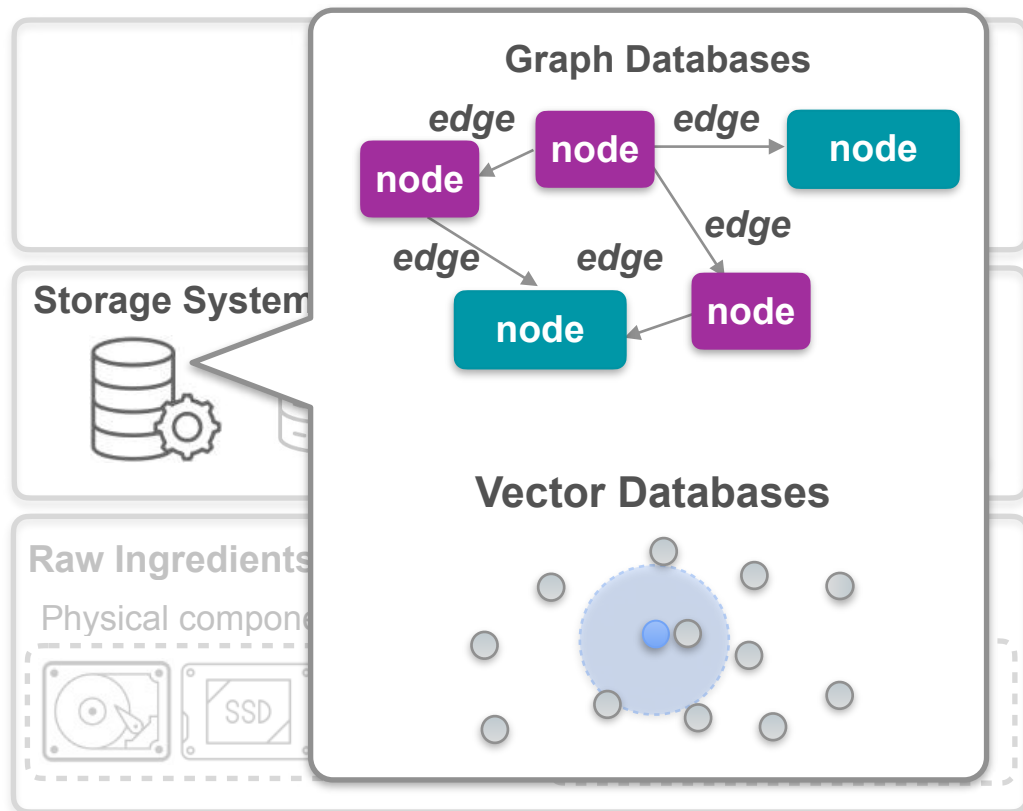
Organizes data in the raw components and allows you to interact with the stored data

OLTP Systems

Online Transactional Processing Systems
Focus on performing read and write queries with low latency

OLAP Systems

Online Analytical Processing Systems
Focus on applying analytical activities on data (e.g. aggregation, summarization)



Storage Hierarchy

Storage Abstractions



Storage Systems



Raw Ingredients

Physical components



Processes

Networking

Serialization

CPU

Compression

Caching

Course 3 Week 1



Trade-offs between storage cost and performance

- Cloud storage paradigms (block, object and file storage)
- Data storage in databases
 - Row vs column-oriented databases
 - Graph and vector databases
- Characteristics of physical components
- Serialization and compression

Storage Abstractions




Storage Systems



Raw Ingredients

Physical components



Processes

Networking

Serialization

CPU

Compression

Caching



Course 3 Week 2

How to choose the appropriate abstractions for storing your data

Storage Abstractions




Storage Systems



Raw Ingredients

Physical components



Processes

Networking

Serialization

CPU

Compression

Caching

Course 3 Week 3



Queries

- How queries work
- How different storage solutions affect query performance
- Techniques for improving query performance

Storage Abstractions




Storage Systems



Raw Ingredients

Physical components



Processes

Networking

Serialization

CPU

Compression

Caching



DeepLearning.AI

Storage Systems

**Raw Ingredients:
Physical Components of Data Storage**

Raw Storage Ingredients

Persistent Storage Medium

Magnetic disk




Solid-state storage



Volatile Memory


RAM



CPU cache




Magnetic Disks



Hard Disk Drives (HDDs)



Magnetic Disks



**Track + Sector
= Address**

Write:


Encode binary data by changing the magnetic field

Read:

Converts magnetic field into binary data

Hard Disk Drives (HDDs)

Solid State Drives





Solid-State Drives (SSDs)

SSDs read and write data
much faster

Performance Comparison

	Magnetic Disk	SSD
Latency	4 milliseconds	
IOPS (Input/output operations per second)	Hundreds	

Commercial magnetic disk drive:
Rotates at 7200 revs/min



Performance Comparison



	Magnetic Disk	SSD
Latency	4 milliseconds	0.1 milliseconds
IOPS (Input/output operations per second)	Hundreds	Tens of thousands



Solid-State Drives (SSDs)




Electrical charges

Performance Comparison



	Magnetic Disk	SSD
Latency	4 milliseconds	0.1 milliseconds
IOPS (Input/output operations per second)	Hundreds	Tens of thousands
Data Transfer Speed (number of bytes read/ written from disk to memory in a second)	Up to 300 MB/s	4 GB/s


Improving Performance



Data transfer speed limited by network performance


Improving Performance

Distributed Storage



Data transfer speed limited by network performance

Partitioning



Slicing SSDs into partitions

Performance Comparison


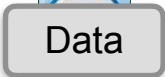





	Magnetic Disk	SSD
Latency	4 milliseconds	0.1 milliseconds
IOPS (Input/output operations per second)	Hundreds	Tens of thousands
Data Transfer Speed (number of bytes read/ written from disk to memory in a second)	Up to 300 MB/s	4 GB/s
Cost	\$0.03–0.06/GB	\$0.08–0.10/GB

2-3 times cheaper

Volatile Memory Ingredients

***Note: these metrics can vary*

	  Magnetic Disk	 SSD	 RAM (Random Access Memory)	 CPU Cache
Latency	4 milliseconds	0.1 milliseconds	0.1 microseconds	
IOPS (Input/output operations per second)	Hundreds	Tens of thousands	Millions	
Data Transfer Speed (number of bytes read/written from disk to memory in a second)	Up to 300 MB/s	4 GB/s	100 GB/s	
Cost	\$0.03–0.06/GB	\$0.08–0.10/GB	> \$3/GB	

30-50 times more expensive

Volatile Memory Ingredients

***Note: these metrics can vary*




	Magnetic Disk	SSD	RAM (Random Access Memory)	CPU Cache
Latency	4 milliseconds	0.1 milliseconds	0.1 microseconds	1 nanosecond
IOPS (Input/output operations per second)	Hundreds	Tens of thousands	Millions	/
Data Transfer Speed (number of bytes read/ written from disk to memory in a second)	Up to 300 MB/s	4 GB/s	100 GB/s	1 TB/s
Cost	\$0.03–0.06/GB	\$0.08–0.10/GB	> \$3/GB	/

CPU Cache Use Cases

- CPU caching
- Store frequently and recently accessed data in a fast access layer

Examples






DeepLearning.AI

Storage Systems

**Raw Ingredients:
Processes Required for Data Storage**

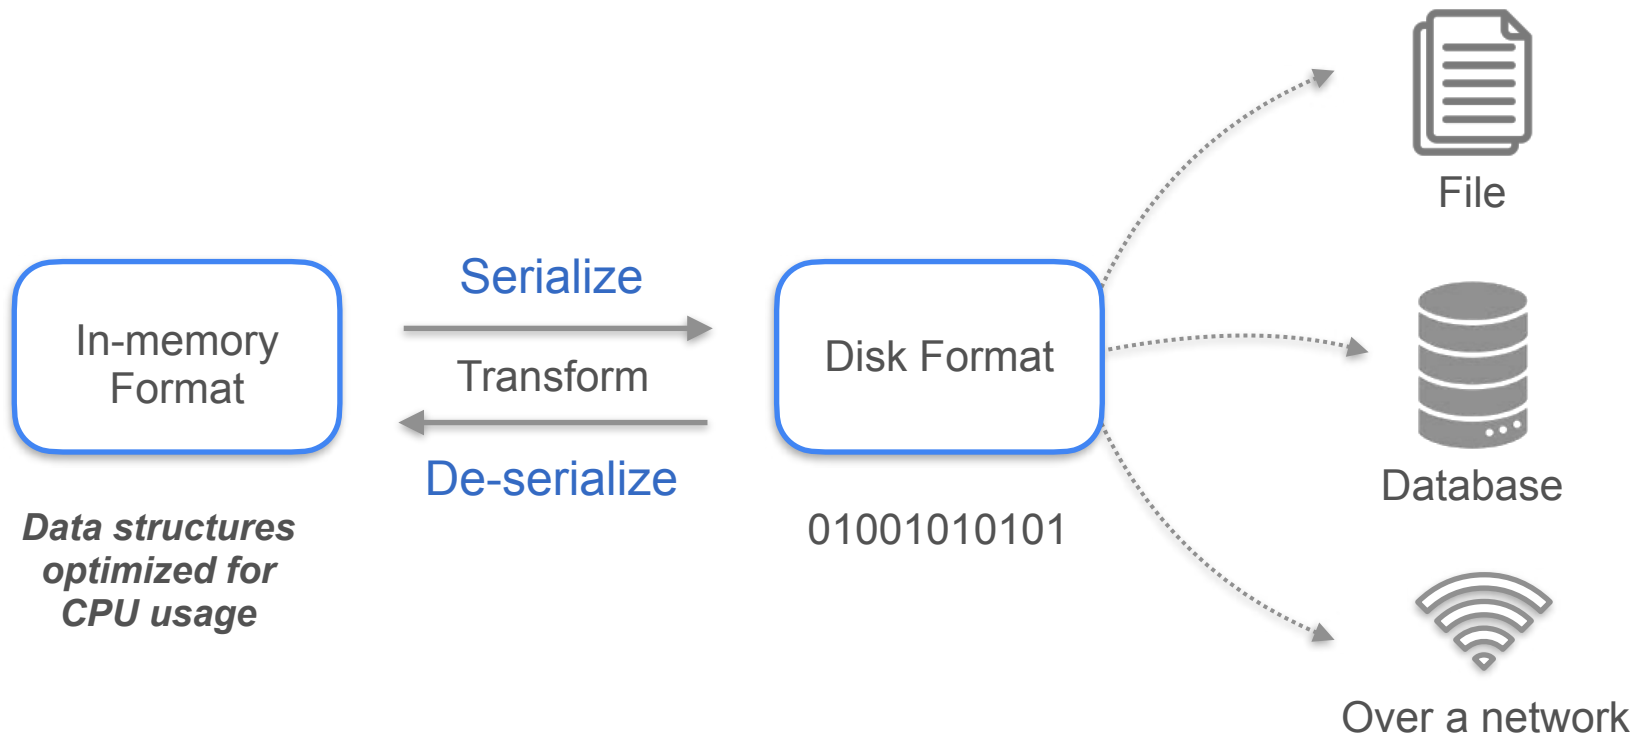
Networking and CPU — “Raw Ingredients” of Storage Systems



Enhance:

- read and write performance
- data durability
- data availability

Serialization



Serialization



Transactional operations

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q



Physical Storage



Row-Based Serialization

bytes representing the 1st object	bytes representing the 2nd object	...	bytes representing the last object
-----------------------------------	-----------------------------------	-----	------------------------------------

Column-Based Serialization

--	--	--	--

Serialization

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

Analytical queries



Physical Storage



Row-Based Serialization

bytes representing the 1st row	bytes representing the 2nd row	...	bytes representing the last row
--------------------------------	--------------------------------	-----	---------------------------------

Column-Based Serialization

bytes representing the 1st	key	bytes representing the 2nd	key	...	bytes representing the last	key
----------------------------	-----	----------------------------	-----	-----	-----------------------------	-----

Serialization Formats

Human-Readable Textual Formats



Binary Formats



Serialization Formats

Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



Binary Formats



Serialization Formats

Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



Binary Formats



Serialization Formats

Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

Binary Formats



Serialization Formats

Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

Binary Formats



- Column-based format
- For efficient storage and big data processing



Serialization Formats

Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

Binary Formats



- Column-based format
- For efficient storage and big data processing



- Row-based format
- Uses a schema to define its data structure
- Supports schema evolution

Serialization Formats

Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

Binary Formats

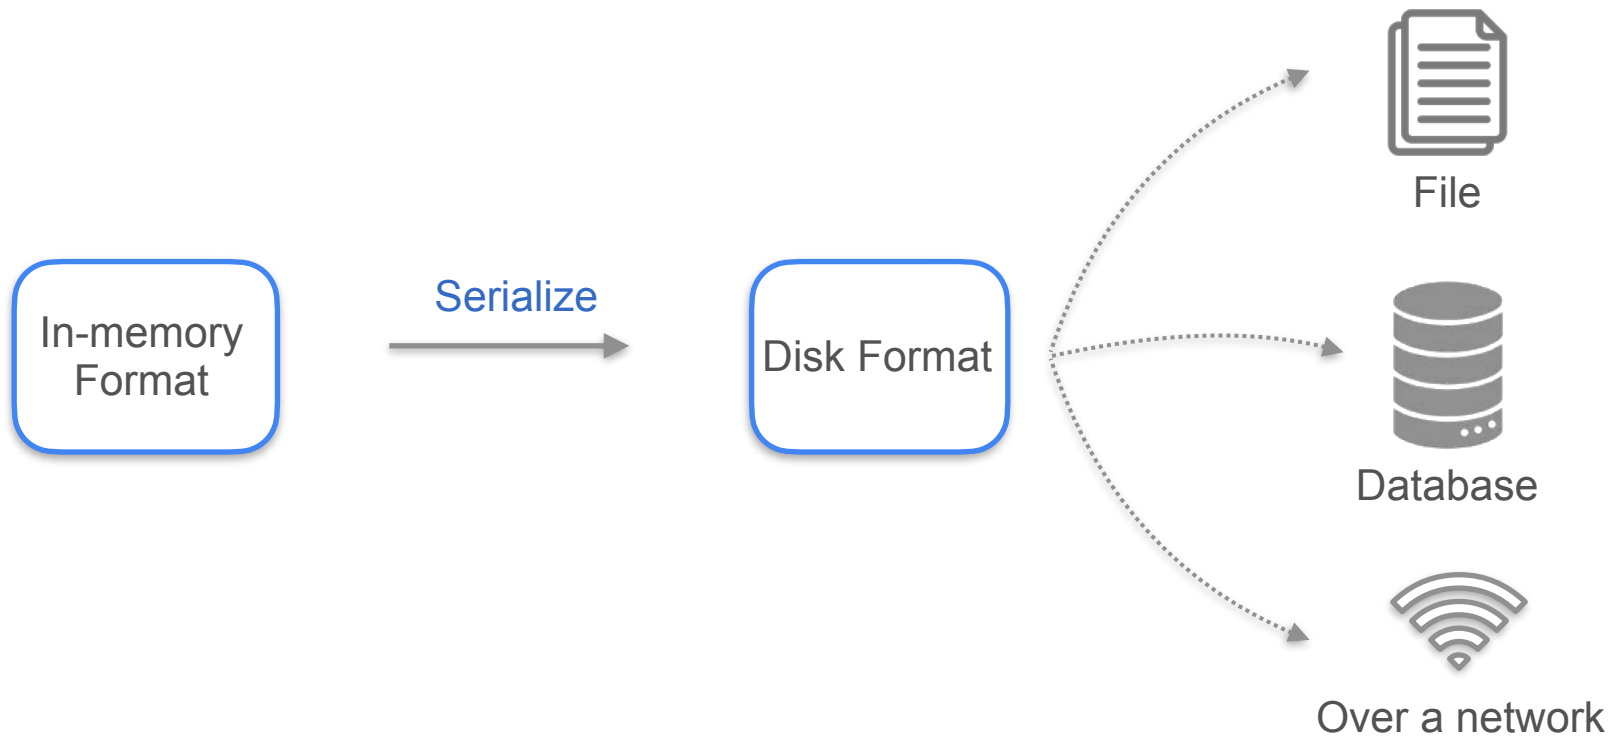


- Column-based format
- For efficient storage and big data processing

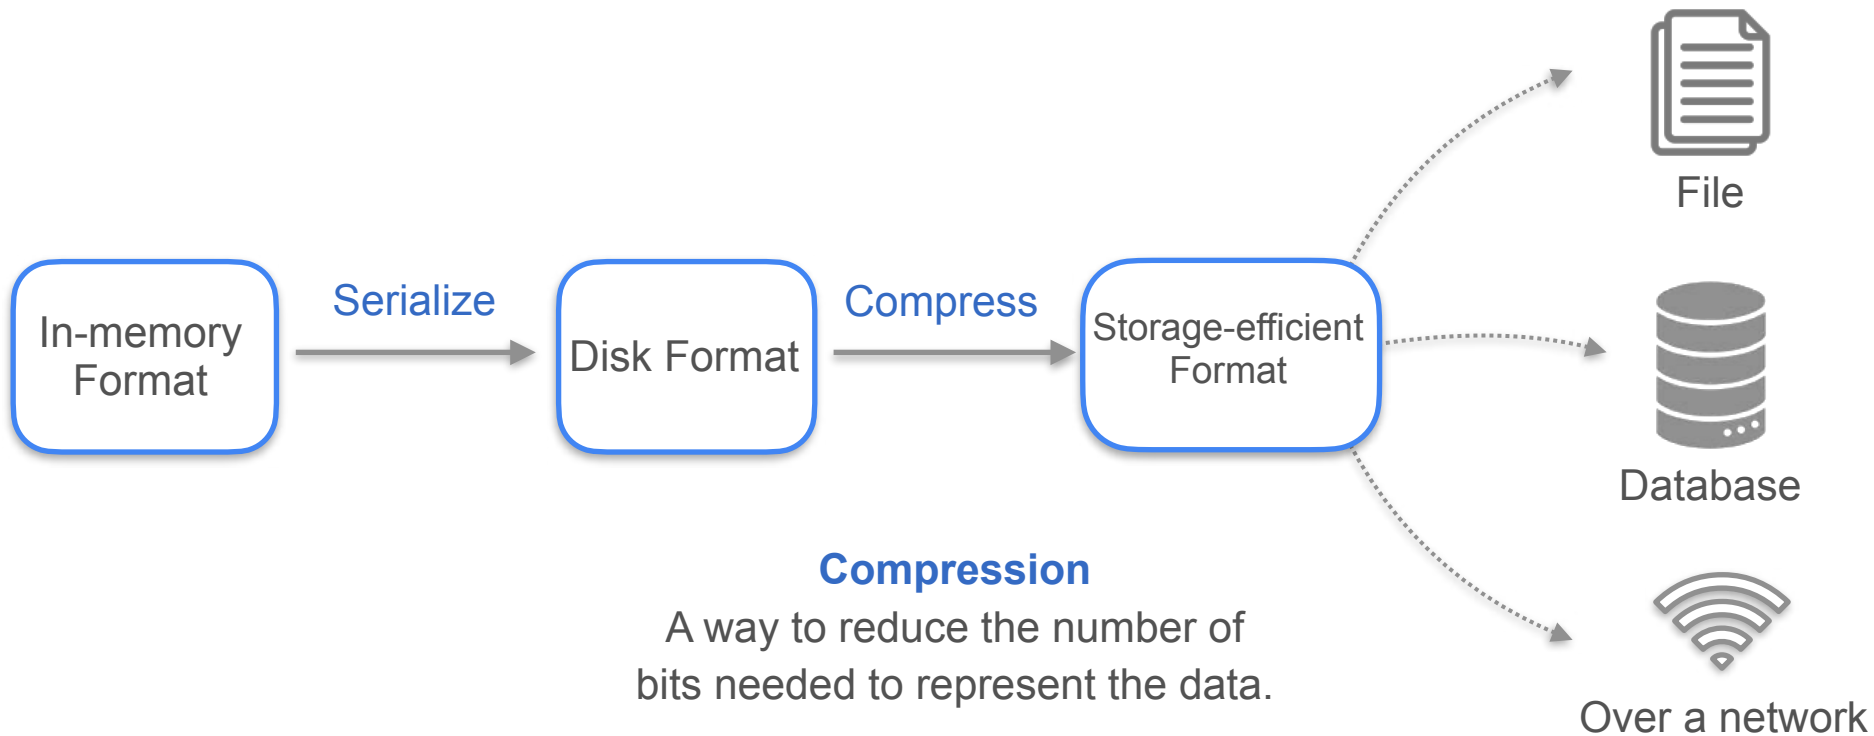


- Row-based format
- Uses a schema to define its data structure
- Supports schema evolution

Serialization

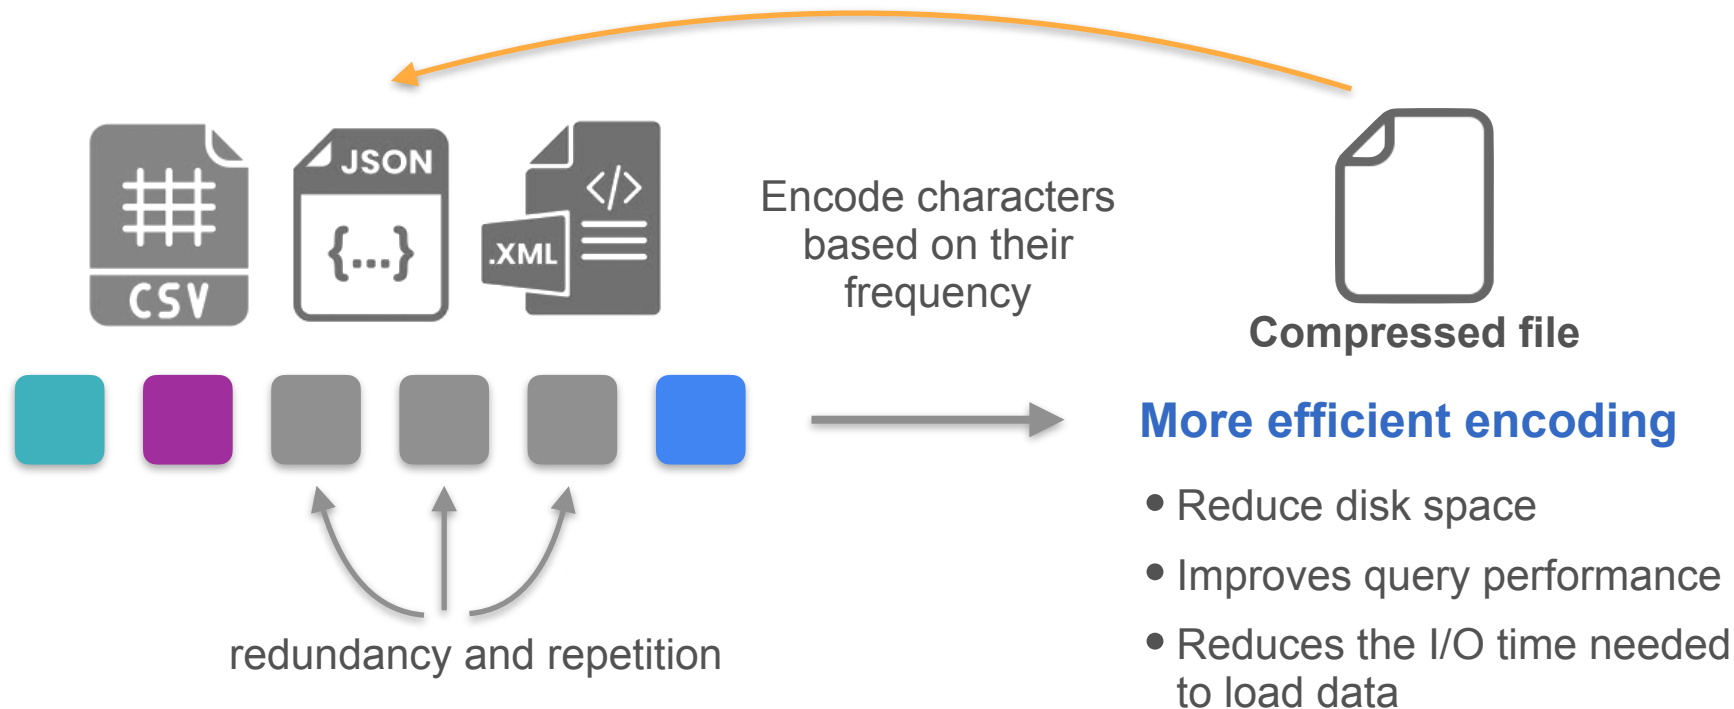


Serialization



Compression

Compression ratio



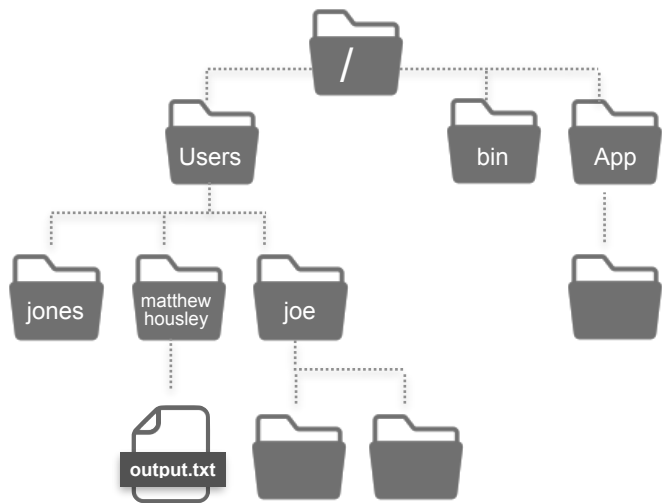


DeepLearning.AI

Storage Systems

**Cloud Storage Options:
Block, Object and File storage**

File Storage




File Storage

Organizes files into a directory tree

Each directory contains metadata about its files and subfolders :

- Name
- Owner
- Last modified date
- Permissions
- Pointer to the actual entity

File Storage



/Users/matthewhousley/output.txt


File Storage

Organizes files into a directory tree

Each directory contains metadata about its files and subfolders :

- Name
- Owner
- Last modified date
- Permissions
- Pointer to the actual entity

File Storage Use Cases




Cloud File Storage Service




Amazon Elastic File System
(EFS)

- Provides you access to shared files over a network
- Networking, scaling, and configuration are handled by the cloud vendor

Block Storage



stored on




Block Storage

Divides files into small, fixed-size blocks of data and stores them on disk

- Each block has a unique identifier
- You can efficiently retrieve and modify data in individual blocks
- You can distribute blocks of data across multiple storage disks
 - Higher scalability
 - Stronger data durability

Block Storage




Lookup Table

File Piece	Block Identifier
First piece	1232
Second piece	1234
Third piece	1236
Fourth piece	1238

Block Storage Use Cases

- Ideal for frequent access and modification
- Enables OLTP systems to perform small and frequent read and write operations with low latency
- Provides persistent storage for virtual machines



Attach a root storage device backed by a block storage volume

Default storage for EC2



Amazon Elastic Block Store (EBS)

1. SSD for latency-sensitive workloads
2. Magnetic disks to store infrequently-accessed data

Object Storage

Object Storage


Stores immutable files as data objects in a flat structure



Object Storage

Object Storage

Stores immutable files as data objects in a flat structure




s3://o'reilly-data-engineering-book/data-example.json

The Bucket

Object key

- Each object is immutable
- To update the file you have to re-write the entire object
- Can scale horizontally and support performant parallel operations



Object Storage Use Cases



Ideal for...	Not ideal for...
<ul style="list-style-type: none">• Storage layer of cloud data warehouses or data lakes• Storing data needed in OLAP systems• Machine learning pipelines<ul style="list-style-type: none">• Raw text• Images• Videos• Audio	<ul style="list-style-type: none">• Not good at supporting transactional workloads



Cloud Storage Options

File Storage	Block Storage	Object Storage
<ul style="list-style-type: none">• Supports data sharing• Easy to manage with low performance and scalability requirements	<ul style="list-style-type: none">• Supports transactional workloads• Allows frequent read and write operations with low latency	<ul style="list-style-type: none">• Supports analytical queries on massive datasets• Offers high scalability and parallel data processing



DeepLearning.AI

Storage Systems

Storage Tiers – Hot, Warm, & Cold Data

Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
Access Frequency			
Example			
Storage Medium			
Storage Cost			
Retrieval Cost			

Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
Access Frequency	Very frequent		
Example	Product recommendation application		
Storage Medium	SSD & Memory		
Storage Cost	High		
Retrieval Cost	Low		

Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
Access Frequency	Very frequent	Less frequent	
Example	Product recommendation application	Regular reports and analyses	
Storage Medium	SSD & Memory	Magnetic disks or hybrid storage systems	
Storage Cost	High	Medium	
Retrieval Cost	Low	Medium	

Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
Access Frequency	Very frequent	Less frequent	Infrequent
Example	Product recommendation application	Regular reports and analyses	Archive
Storage Medium	SSD & Memory	Magnetic disks or hybrid storage systems	Low-cost magnetic disks
Storage Cost	High	Medium	Low
Retrieval Cost	Low	Medium	High

Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
Access Frequency	Very frequent	Less frequent	Infrequent
Example	Product recommendation application	Regular reports and analyses	Archive
Storage Medium	SSD & Memory	Magnetic disks or hybrid storage systems	Low-cost magnetic disks
Storage Cost	High	Medium	Low
Retrieval Cost	Low	Medium	High



AWS Storage Tiers



Amazon S3

Access Frequency

Hot Storage



S3 Express
One Zone



S3 Standard



S3 Standard-
IA



S3 One
Zone-IA



S3 Glacier
Flexible
Retrieval



S3 Glacier
Deep
Archive



S3 Glacier
Instant
Retrieval




DeepLearning.AI


Storage Systems

Distributed Storage Systems

How Distributed Storage Systems Work




How Distributed Storage Systems Work



Magnetic disk




SSD




Total Capacity = + +


How Distributed Storage Systems Work



Single Machine Storage Architecture




How Distributed Storage Systems Work




- Higher fault tolerance and data durability
- High availability

How Distributed Storage Systems Work



- Higher fault tolerance and data durability
- High availability
- Process many read and write operations in parallel
- Fast data access



Advantages of Distributed Storage Systems

Distributed Storage Architecture



Object Storage



Cloud
Data Warehouse



Methods for Distributing Data

Replication

Partitioning

Methods for Distributing Data



Replication



High availability and performance

Partitioning

Methods for Distributing Data

Replication



High availability and performance

Partitioning



Sharding



← partition or shard


Methods for Distributing Data

Replication



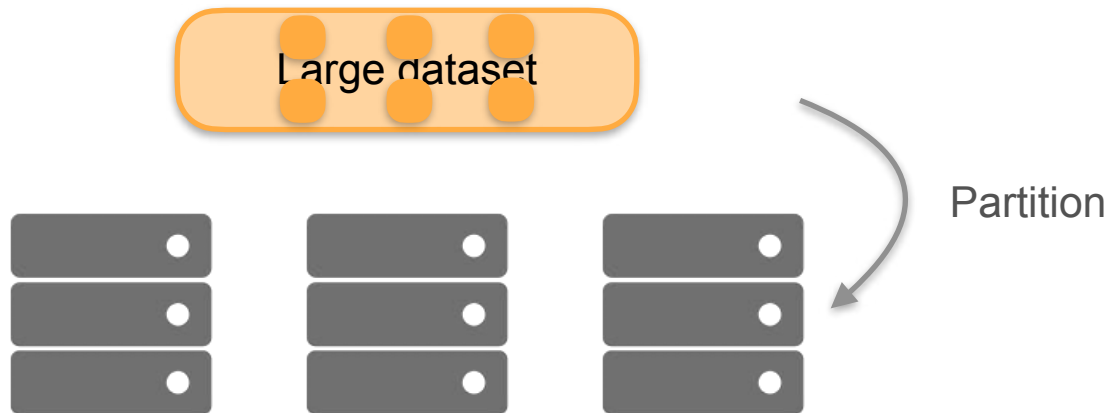
High availability and performance

Partitioning

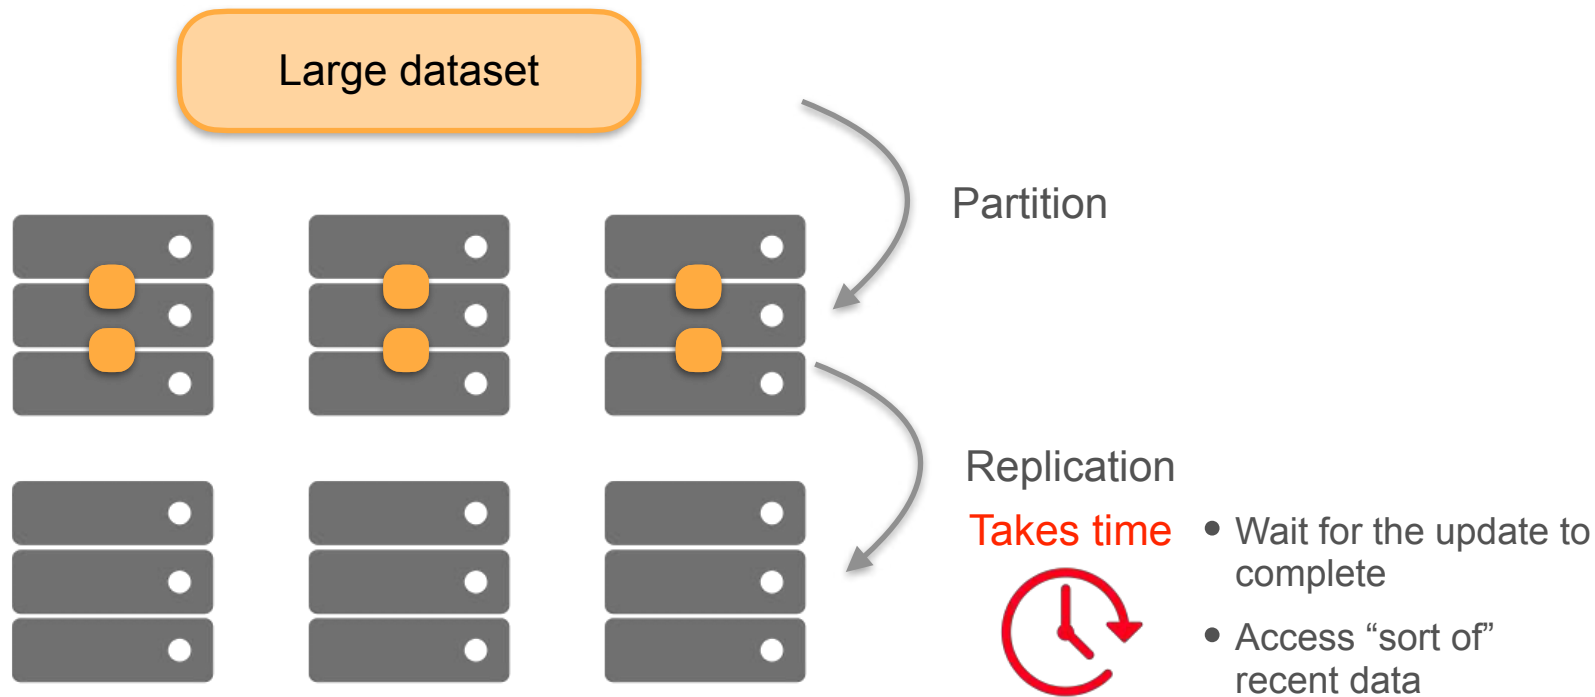


Sharding

Methods for Distributing Data




Methods for Distributing Data



Distributed Storage Considerations – CAP Theorem

The CAP theorem



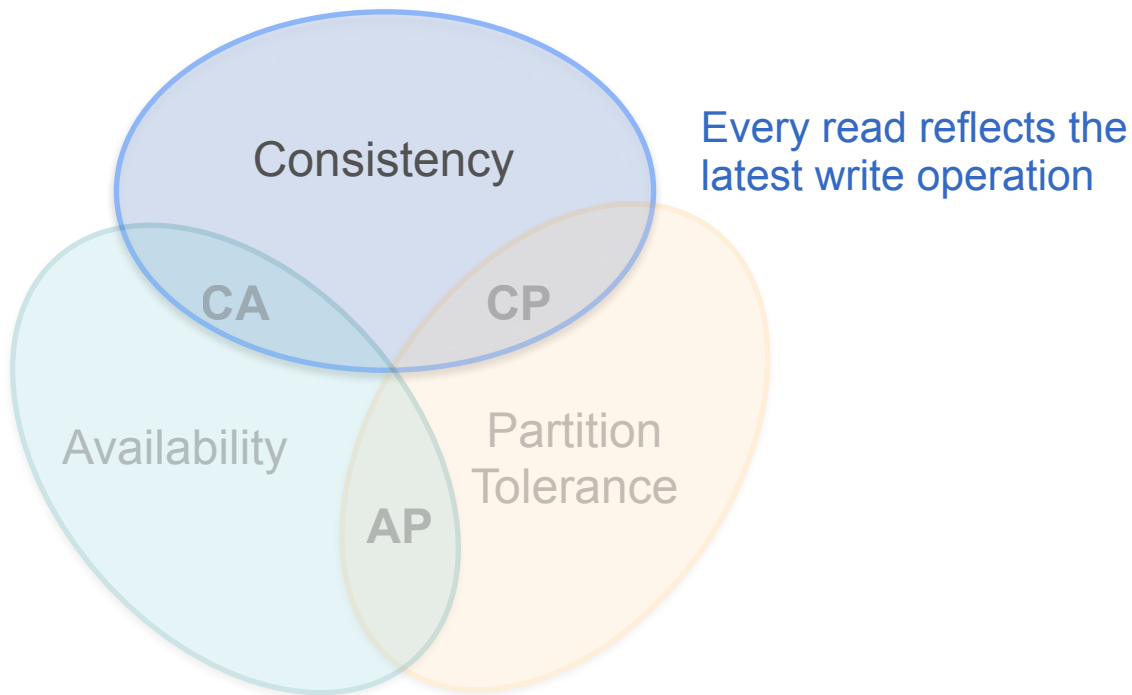
Distributed Storage Considerations – CAP Theorem

ACID




Consistency

Any change to data must follow the set of rules defined by database schema

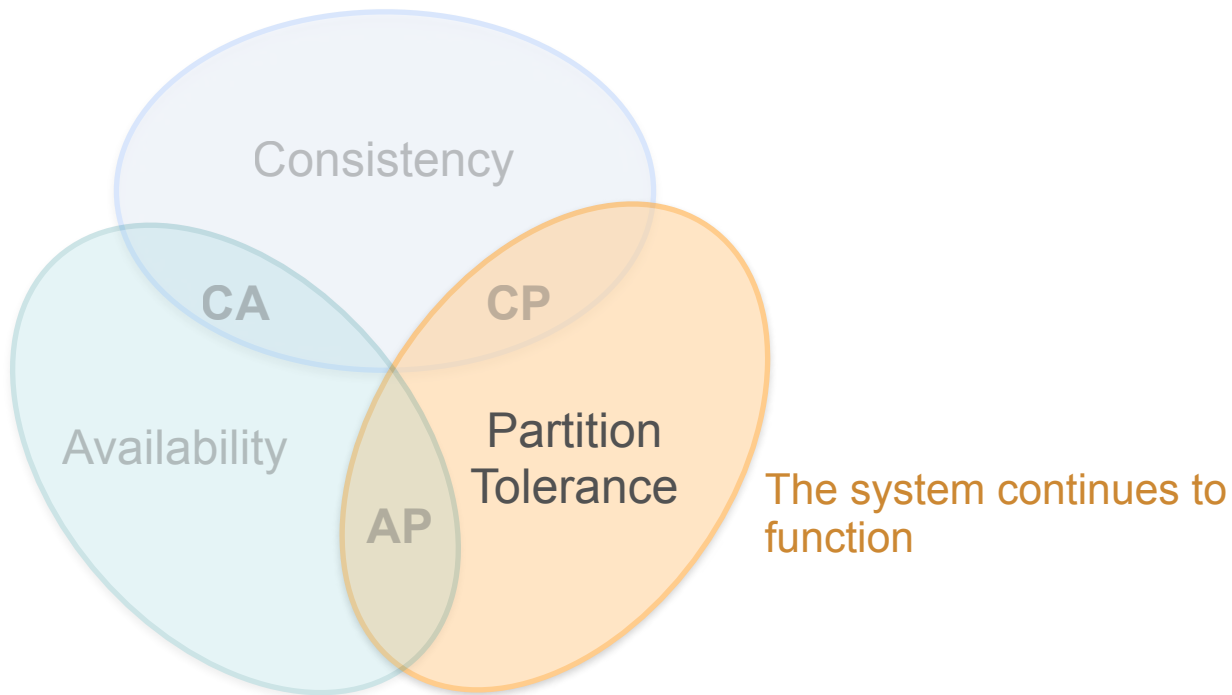


Distributed Storage Considerations – CAP Theorem



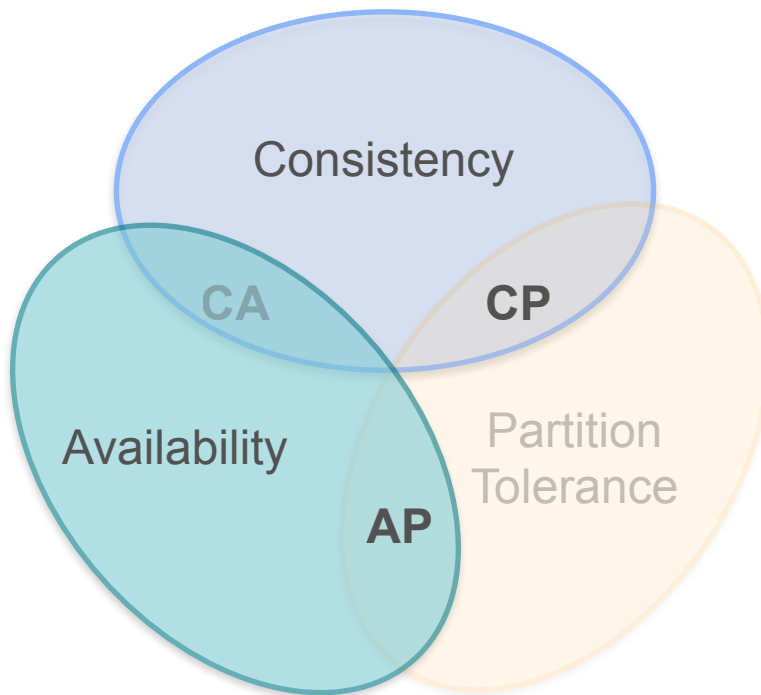
Every request will receive a response

Distributed Storage Considerations – CAP Theorem



Distributed Storage Considerations – CAP Theorem

The CAP theorem



Distributed Storage Considerations – CAP Theorem

Scenario:


Accessing a node that's still being updated

Option 1:

Cancel the request


Option 2:

Proceed with the read operation



Distributed storage considerations – ACID vs BASE

RDBMS




ACID
compliant

NoSQL Databases



BASE
principles



Distributed storage considerations – ACID vs BASE

ACID

compliant

Atomicity

Consistency

Isolation

Durability

BASE

principles

Basically **A**vailable

Soft-state

Eventual Consistency

Scenario

Course 1



Data Scientist



Main database instance (Strong consistency)



Read-replica of the prod database

- Ingest
- Transform
- Store
- Serve

Read Replicas in RDS (Eventual consistency)



- Track changes in main database
- Update their own data



Amazon RDS



DeepLearning.AI

Lab Walkthrough

Comparing Cloud Storage Options

Object Storage

File Storage

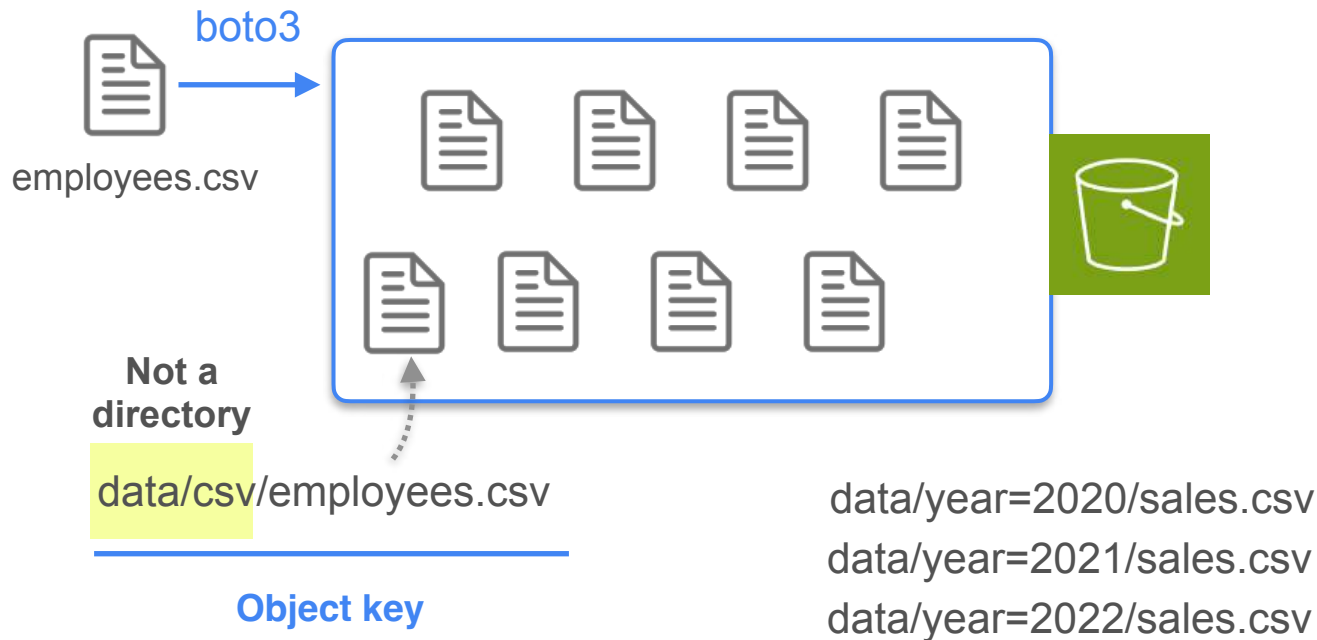
Block Storage

Memory

Object Storage

Flat Structure

Immutability



Object Storage

Flat Structure

Immutability

File Storage

Block Storage

Memory

2. Modify employees' data

 `boto3`
`employees.csv`

Not a
directory

`data/csv/employees.csv`

Object key



1. Enable versioning

3. Use
`list_object_version`

Object Key: data/csv/employees.csv
Object Version Id: q4A5B9CQZ10.u7.YKA7LabC_5GcBA.uZ
Is Latest: True
Last Modified: 2024-08-12 19:57:00+00:00

Object Key: data/csv/employees.csv
Object Version Id: WOOPNaVQTFHB13CkIiRATRDCZt30Wq9k
Is Latest: False
Last Modified: 2024-08-12 19:39:56+00:00

Object Storage

File Storage

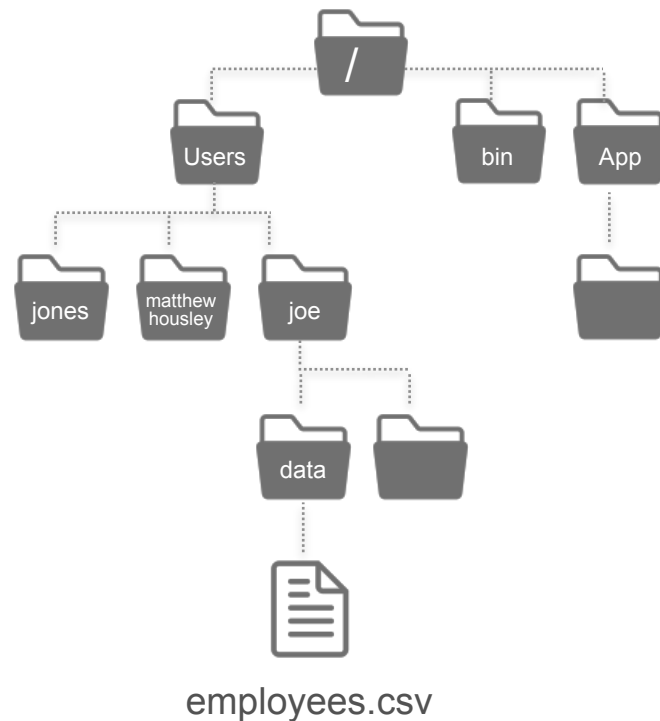
Block Storage

Memory

1. Navigate to the “data” directory
2. Explore the directory content and metadata
3. Explore how the data is modified in place

A directory

data/employees.csv



Object Storage

File Storage

Block Storage

Memory

1. Connect to the server
2. Send a file to the server



Server that emulates the behavior of block storage



Object Storage

Transferring data from memory is faster than transferring data from disk.



Use the `cache-pandas` package:

- provides the “`timed_LRU_cache`” decorator to easily cache in memory pandas DataFrames generated by functions.

```
@timed_lru_cache(seconds=100, maxsize=None)
def read_csv_to_memory(path: str) -> pd.DataFrame:
    """Read CSV function with a cache decorator."""
    return pd.read_csv(path)
```


Block Storage

Memory


Compare the time it takes to read the file for the first time with the time it takes to read the same data stored in memory.

Object Storage

Monitor your memory storage capacity using
htop command:



File Storage



PID	USER	PRI	NI	VRT	RES	SHR	S	CPU%	MEM%	TIME	Command
1	root	20	0	166M	15960	9756	S	0.0	0.8	0:01.59	/usr/lib/systemd/systemd --switched-root --system --deserialize=32
788	root	20	0	52944	17620	16364	S	0.0	0.9	0:00.72	/usr/lib/systemd/systemd-journald
1207	root	20	0	30880	10384	7572	S	0.0	0.5	0:00.07	/usr/lib/systemd/systemd-udev
1387	systemd-re	20	0	21212	13876	10608	S	0.0	0.7	0:00.11	/usr/lib/systemd/systemd-resolved
1392	root	16	-4	20224	2300	1632	S	0.0	0.1	0:00.03	/sbin/auditd
1393	root	16	-4	20224	2300	1632	S	0.0	0.1	0:00.00	/sbin/auditd
1424	dbus	20	0	8380	3932	3344	S	0.0	0.2	0:00.02	/usr/bin/dbus-broker-launch --scope system --audit
1425	dbus	20	0	5264	2904	2404	S	0.0	0.1	0:00.19	dbus-broker --log 4 --controller 9 --machine-id ec20b71fb5d5e6c4dae675b053431827 --max
1426	root	20	0	15296	6444	5612	S	0.0	0.3	0:00.01	/usr/bin/systemd-inhibit --what-handle-suspend-key:handle-hibernate-key --who=noah --w
1428	root	20	0	81344	3024	2796	S	0.0	0.2	0:00.15	/usr/sbin/irqbalance --foreground
1429	libstorage	20	0	2756	2072	1908	S	0.0	0.1	0:00.03	/usr/bin/lsnd -d
1431	root	20	0	161M	6600	5396	S	0.0	0.3	0:37.39	/usr/sbin/rngd -f -x pkcs11 -x nist
1433	root	20	0	15800	7640	6672	S	0.0	0.4	0:00.02	/usr/lib/systemd/systemd-homed
1434	root	20	0	17744	9604	7520	S	0.0	0.5	0:00.16	/usr/lib/systemd/systemd-logind
1436	systemd-ne	20	0	230M	9220	7956	S	0.0	0.5	0:00.08	/usr/lib/systemd/systemd-networkd
1440	root	20	0	81344	3024	2796	S	0.0	0.2	0:00.00	/usr/sbin/irqbalance --foreground
1459	root	20	0	2668	1108	1020	S	0.0	0.1	0:00.00	/usr/sbin/acpid -f
1460	root	20	0	161M	6600	5396	S	0.0	0.3	0:18.47	/usr/sbin/rngd -f -x pkcs11 -x nist
1461	root	20	0	161M	6600	5396	S	0.0	0.3	0:18.87	/usr/sbin/rngd -f -x pkcs11 -x nist
1468	root	20	0	274M	3396	2632	S	0.0	0.2	0:00.00	/usr/sbin/gssproxy -D
1470	root	20	0	274M	3396	2632	S	0.0	0.2	0:00.00	/usr/sbin/gssproxy -D
1471	root	20	0	274M	3396	2632	S	0.0	0.2	0:00.00	/usr/sbin/gssproxy -D
1472	root	20	0	274M	3396	2632	S	0.0	0.2	0:00.00	/usr/sbin/gssproxy -D
1473	root	20	0	274M	3396	2632	S	0.0	0.2	0:00.00	/usr/sbin/gssproxy -D
1474	root	20	0	274M	3396	2632	S	0.0	0.2	0:00.00	/usr/sbin/gssproxy -D
1475	root	20	0	1776M	20824	8204	S	0.0	1.1	0:05.19	/usr/bin/containerd
1516	root	20	0	1776M	20824	8204	S	0.0	1.1	0:01.46	/usr/bin/containerd
1523	root	20	0	1776M	20824	8204	S	0.0	1.1	0:00.00	/usr/bin/containerd
1524	root	20	0	1776M	20824	8204	S	0.0	1.1	0:00.46	/usr/bin/containerd
1534	root	20	0	1776M	20824	8204	S	0.0	1.1	0:00.00	/usr/bin/containerd

F1 Help F2 Setup F3 Search F4 Filter F5 Tree F6 Sort F7 Nice F8 Force F9 Kill F10 Quit

Object Storage

File Storage

Block Storage

Memory

Explore the features of these storage options.

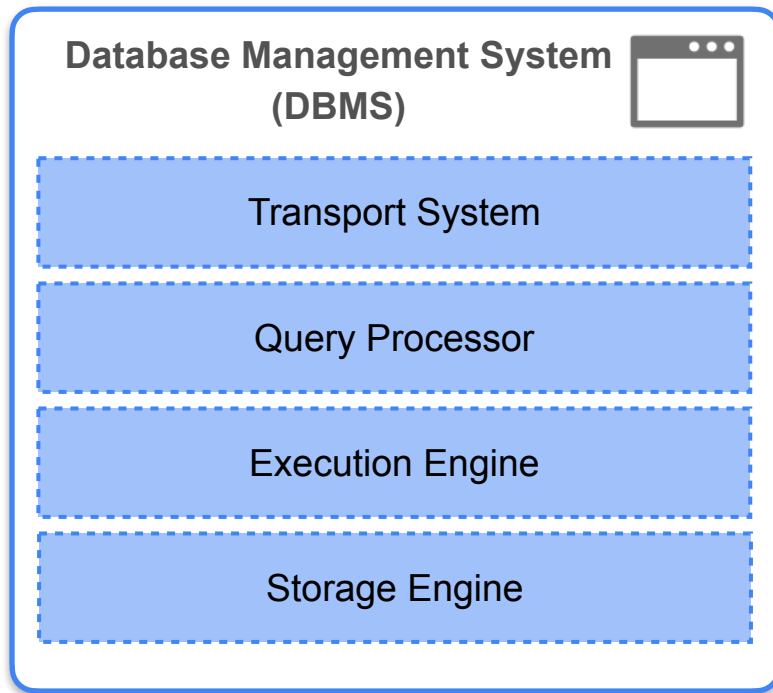


DeepLearning.AI

Storage Systems

How Databases Store Data

Database Management System



Database Management System

Storage Engine

- Serialization
- Arrangement of data on disk
- Indexing

Modern Storage Engines

- Support the performance characteristics of SSDs
- Handle modern data types and structures
- Offer robust columnar storage support

Average price of products purchased in the USA

```
SELECT AVG(price)
```

```
FROM my_table
```

```
WHERE country = "USA"
```

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico

.....

Average price of products purchased in the USA

```
SELECT AVG(price)
FROM my_table
WHERE country = "USA"
```

Index

A data structure that helps you efficiently locate data

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico

.....

scan all rows

Average price of products purchased in the USA

```
SELECT AVG(price)
FROM my_table
WHERE country = "USA"
```

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico

***Scanning all rows: $O(n)$

Binary search on rows: $O(\log n)$

Index


A data structure that helps you efficiently locate data

Index table

Country	Row Address
Brazil	###
Canada	###
Canada	###
Canada	###
Canada	###
Canada	###
Mexico	###
USA	###
USA	###
USA	###

Use binary search to locate the USA rows

In-Memory Storage Systems



- Excellent transfer speed and low latency
- Volatile
- Used to present data for ultra-fast retrieval:
 - Caching applications
 - Real-time bidding
 - Gaming leaderboards

1. Memcached

- Key-value store to cache database query results or API calls
- Used when it's acceptable for data to be lost

2. Redis

- Key-value store that supports more complex data types
- Supports high-performance applications that can tolerate minor data loss



DeepLearning.AI

Storage Systems

Row vs Column Storage

Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

↓ Stores data
row by row

Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
bytes representing the 1st row					bytes representing the 2nd row					bytes representing the last row					

Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

Row Storage is perfect for OLTP

Perform read and write operations with low latency

← Locate this order

↓
Stores data
row by row

Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----

Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q



Stores data
row by row

Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----

Analytical queries focus on summarizing or aggregating columns

- Total revenue?
- Most popular product?
- Average quantity?

Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...

```
SELECT SUM(price)
FROM my_table
```

1 million rows

30 columns

100 bytes per entry

Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q	...
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----	-----

Row-Oriented Storage


Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...

1 million rows

30 columns

100 bytes per entry

```
SELECT SUM(price)
FROM my_table
```



CPU

Physical Storage



Row-Oriented Storage


Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...

1 million rows \times 30 columns \times 100 bytes per entry = 3 GB

Data transfer speed: 200 MB/s

Total transfer time? $\frac{3\text{GB or } 3000\text{ MB}}{200\text{ MB/s}} = 15\text{ s}$

```
SELECT SUM(price)
FROM my_table
```



Row-Oriented Storage


Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...

1 billion rows X 30 columns X 100 bytes per entry = 3000 GB

Data transfer speed: 200 MB/s

Total transfer time? $\frac{3000 \text{ GB}}{200 \text{ MB/s}} = 4 \text{ hours !}$

```
SELECT SUM(price)
FROM my_table
```



Column-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...

↓ Stores data
Column by column

Physical Storage

1	2	3	4	40	23	45	50	45865	90234	12558	45682	...
---	---	---	---	----	----	----	----	-------	-------	-------	-------	-----

bytes representing 1st column bytes representing 2nd column bytes representing 3rd column

Column-Oriented Storage — Suitable for OLAP systems!


Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...

1 billion rows 30 columns 100 bytes per entry = 100 GB

Data transfer speed: 200 MB/s

Total transfer time? $\frac{100 \text{ GB or } 100,000 \text{ MB}}{200 \text{ MB/s}} = 8.33 \text{ minutes}$

```
SELECT SUM(price)
FROM my_table
```



Row-oriented Storage

Transfer 1 billion rows
from disk to memory

4 hours

Column-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...

Terrible for
transactional
workloads!

Stores data
Column by column

Physical Storage

1	2	3	4	40	23	45	50	45865	90234	12558	45682	...
---	---	---	---	----	----	----	----	-------	-------	-------	-------	-----

Deserialize the column, modify it,
then write it back to storage




DeepLearning.AI

Storage Systems

Graph Databases

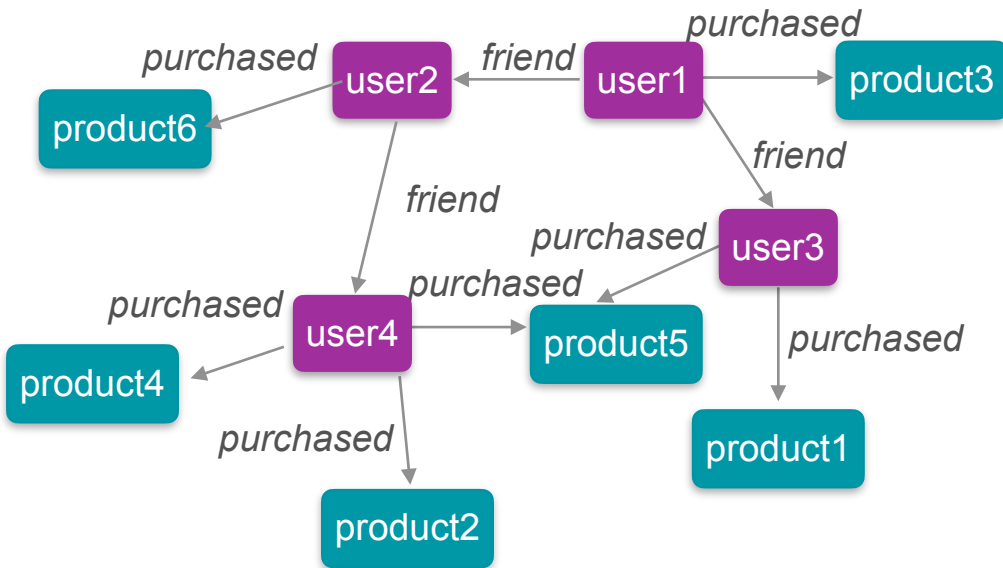
Graph Database



- Nodes represent data items
- Edges represent connection between the data items
- Graph databases model complex connections between data entities

Graph Database

Relationships are first-class citizens



Relational database

purchase

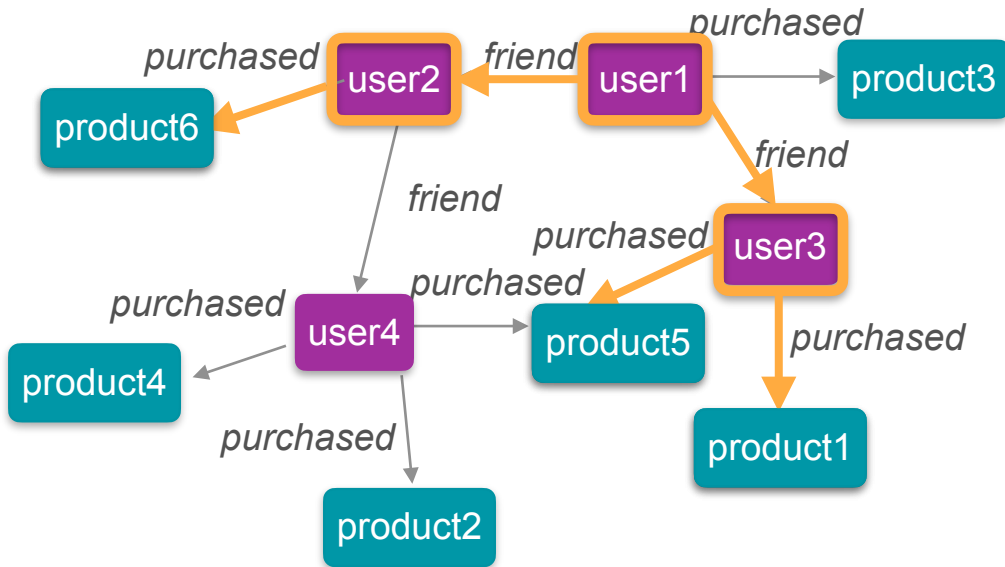
user	product
user1	product3
user2	product6
user3	product1
user3	product5
user4	product5
user4	product4
user4	product2

friendship

user	friend
user1	user3
user1	user2
user2	user4

Querying Data

Traverse the graph to query relationships



Recommendations
for user1

Relational database

purchase

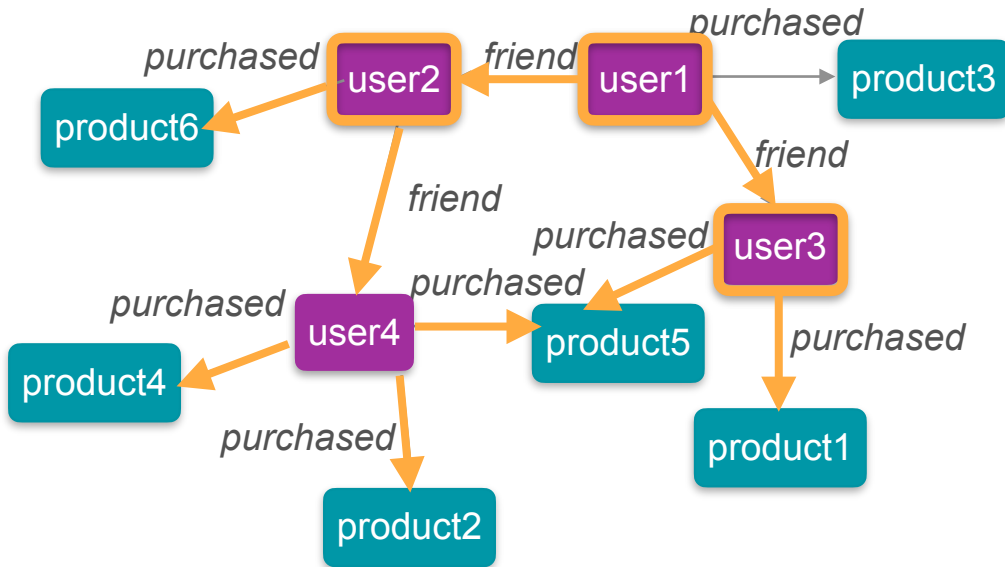
user	product
user1	product3
user2	product6
user3	product1
user3	product5
user4	product5
user4	product4
user4	product2

friendship

user	friend
user1	user3
user1	user2
user2	user4

Querying Data

Traverse the graph to query relationships



Recommendations
for user1

product1

product5

product6

Relational database

purchase

user	product
user1	product3
user2	product6
user3	product1
user3	product5
user4	product5
user4	product4
user4	product2

friendship

user	friend
user1	user3
user1	user2
user2	user4

Less efficient in
querying complex
relationships!

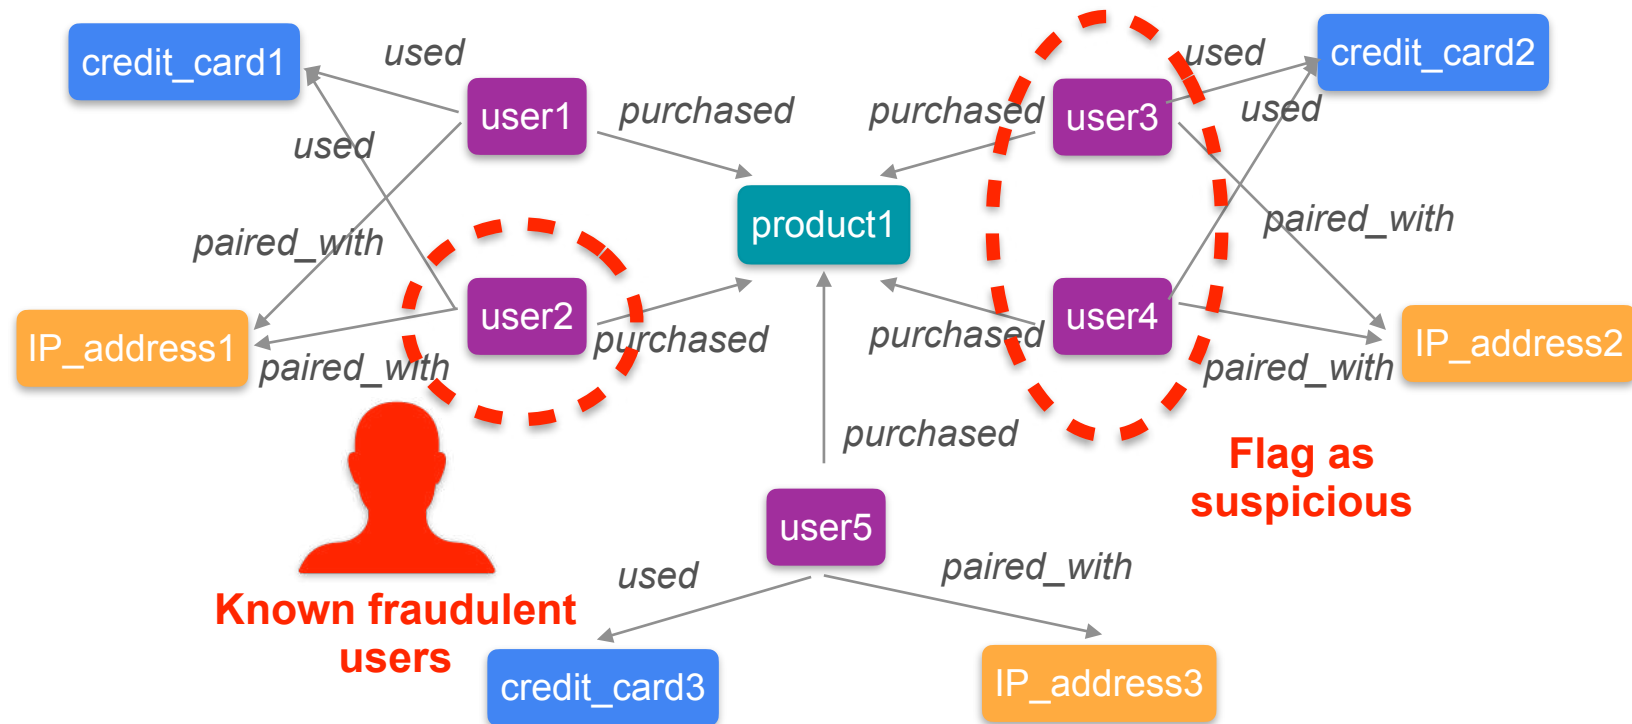
```
SELECT DISTINCT purchase.product
FROM friendship
```

```
JOIN purchase ON friendship.friend = purchase.user
WHERE friendship.user = 'user1'
```

Graph Database - Use Cases

- Recommending products
- Modeling social networks
- Representing network and IT operations
- Simulating supply chains logistics
- Tracing data lineage

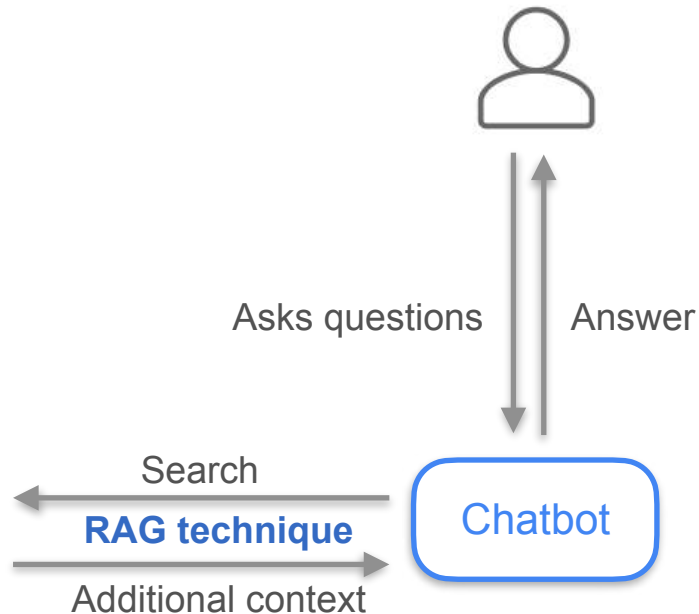

Use Case - Fraud Detection



Use Case - Knowledge Graph

Knowledge Graph

Connect diverse data from disparate sources



Graph Databases

Examples of Graph Databases



Amazon Neptune

Examples of Graph Query Language

Cypher

Gremlin

SparQL



DeepLearning.AI

Storage Systems

Vector Databases

Vector data

Consists of numerical values arranged in an array



Image



0.5	0.5	0.7	0.5	0.3	0.1	0.7	0.2
0.4	0.8	0.9	0.1	0.3	0.1	0.4	0.2
0.3	0.5	0.7	0.8	0.3	0.1	0.6	0.2

Vector embeddings


Capture semantic meaning of an item, like a text document or image




- Can convert an entire database of docs or text into embeddings
- Embeddings help you more efficiently find and retrieve similar items
- Example: Finding similar text
 - Compute embeddings for the query item
 - Database returns similar vectors (based on closeness)

Distance Metric


Vector database uses a distance metric to find similar vectors



Euclidean distance




Cosine distance



Manhattan distance

Similarity Search - Popular Algorithm


K-nearest neighbors (KNN)



- Calculates distance to all vector embeddings
- Becomes inefficient when the data size increases
- Suffers from the curse of dimensionality

ANN (Approximate Nearest Neighbors)

- Find a good guess for the nearest neighbors
- More efficient than K-NN
- Vector databases are built to support ANN algorithms

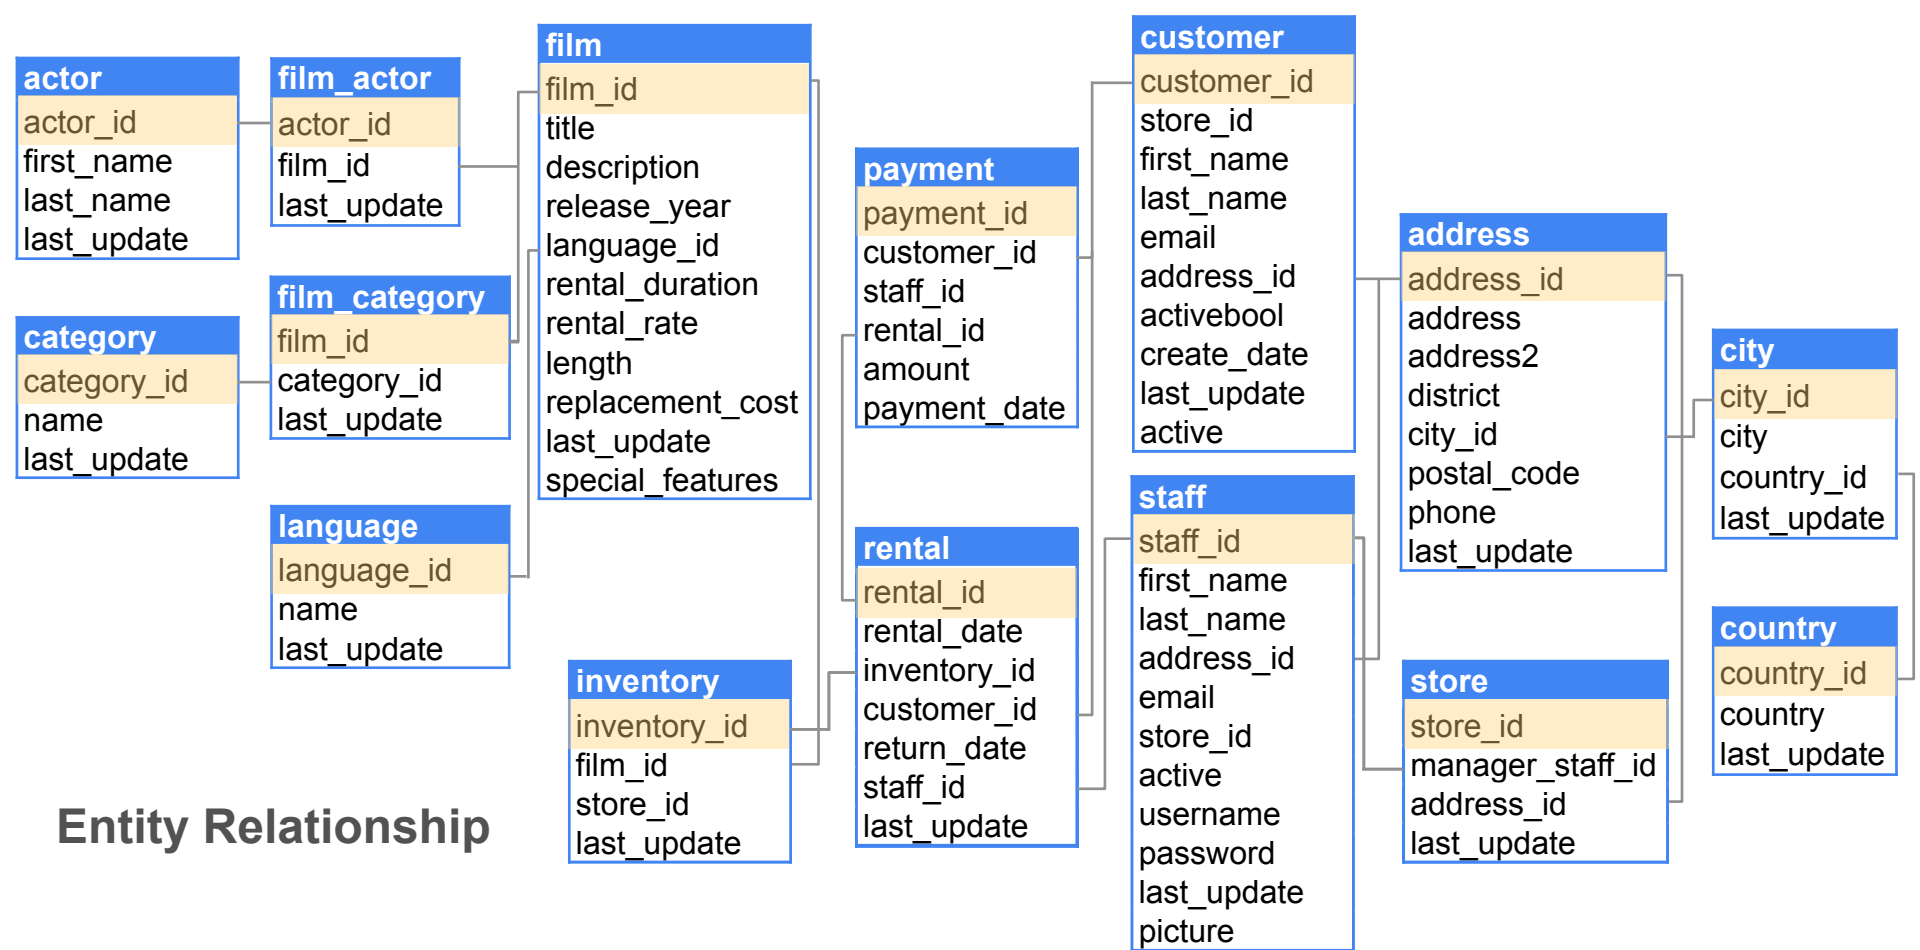




DeepLearning.AI


Storage Systems

Neo4j Graph Database & Cypher Query Language (Part 1)




Entity Relationship

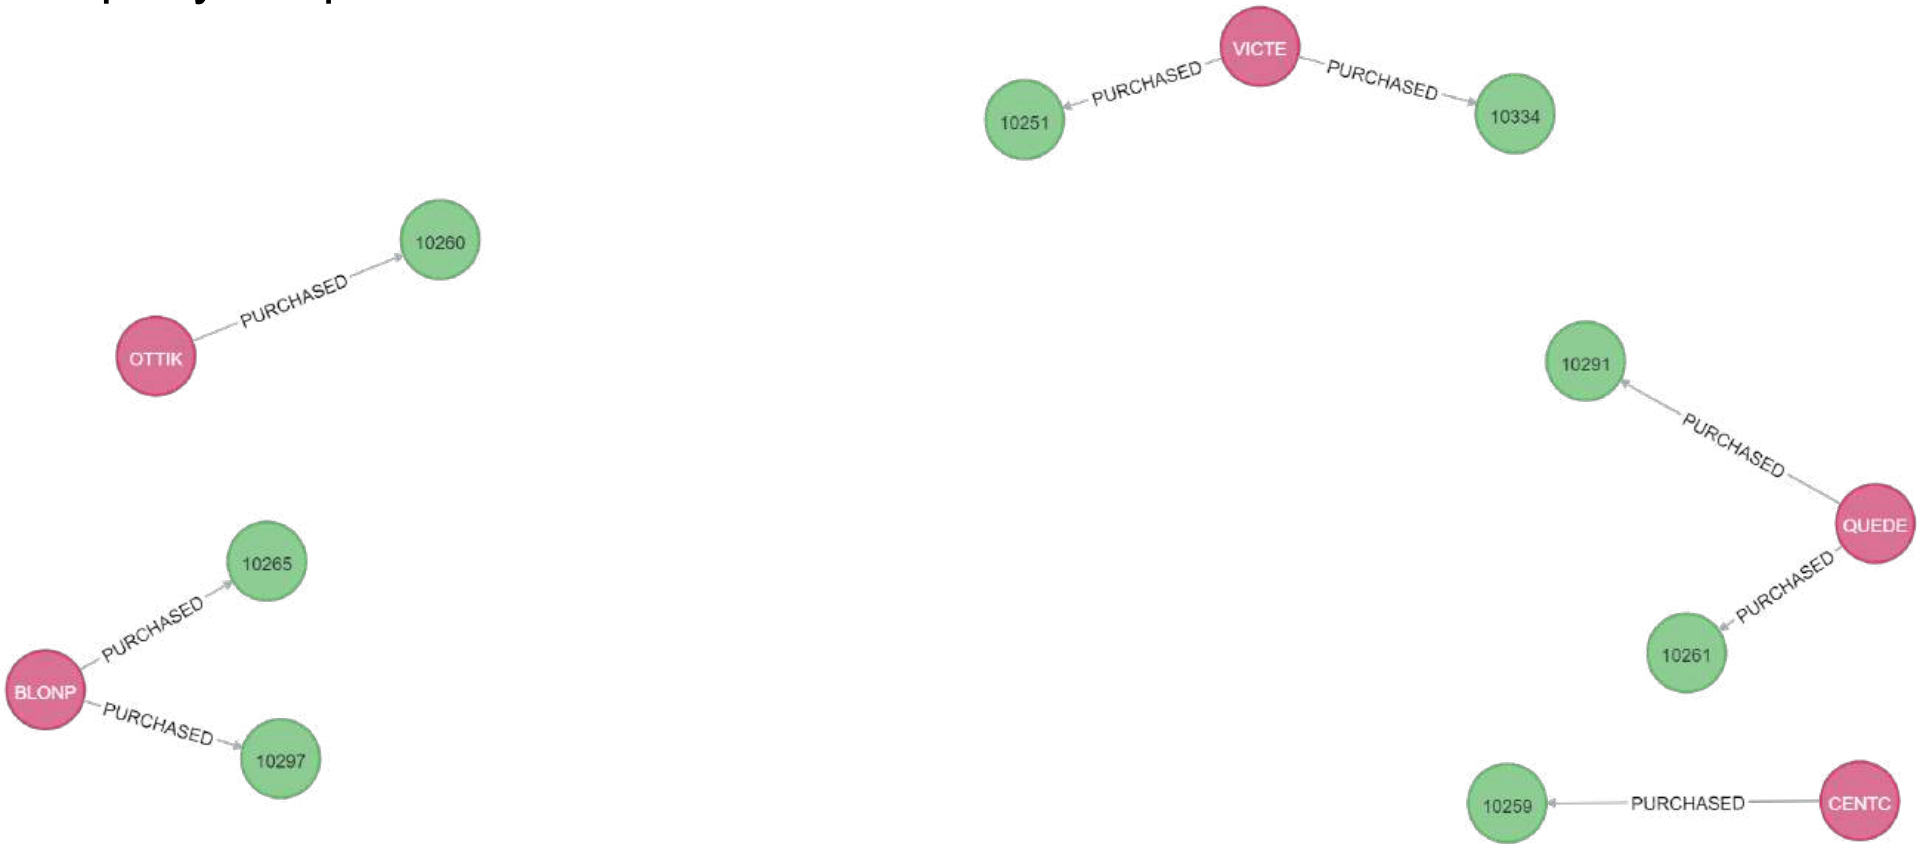
Property Graph Model



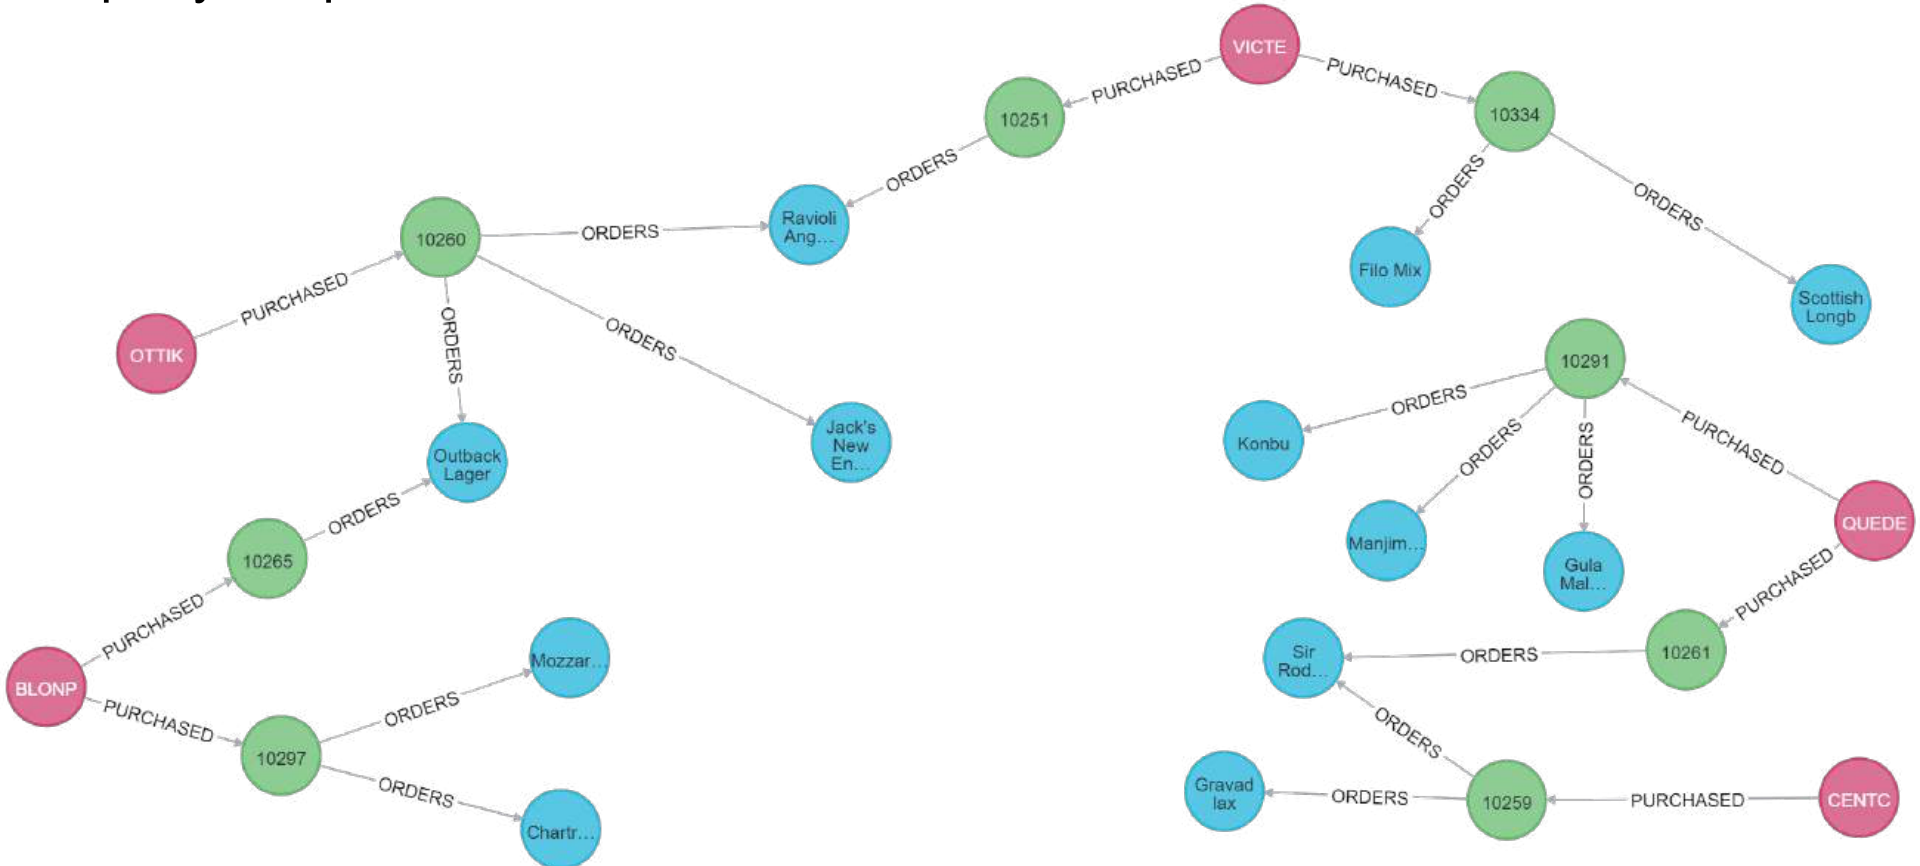
Property Graph Model



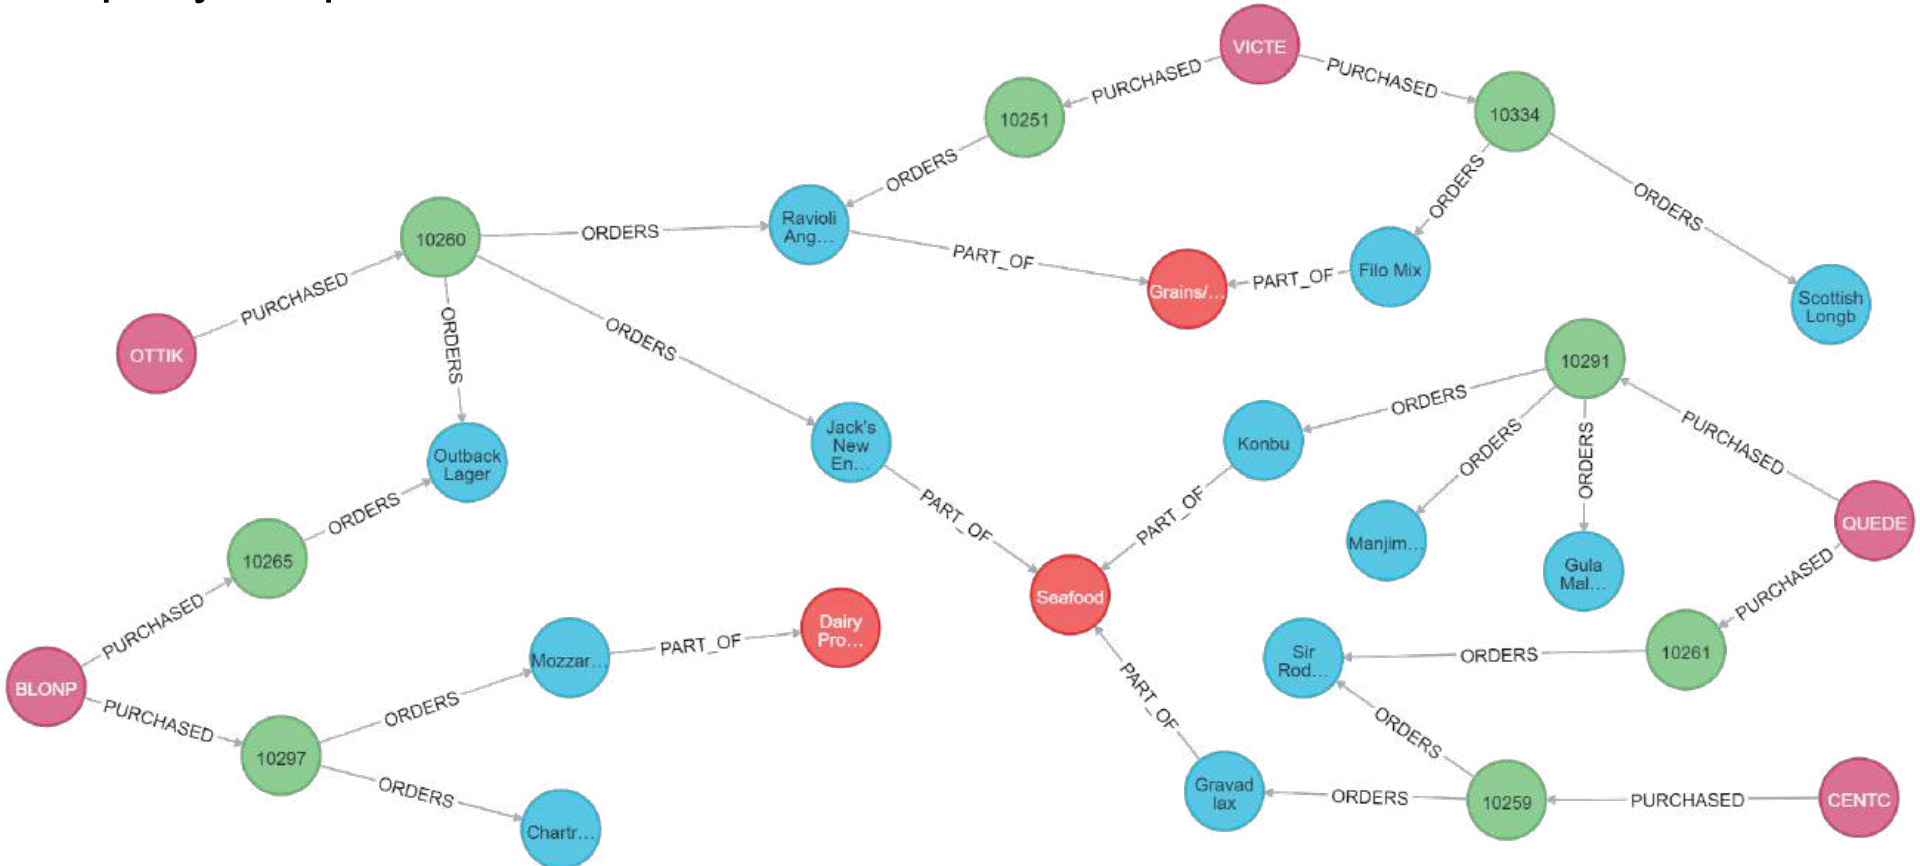
Property Graph Model



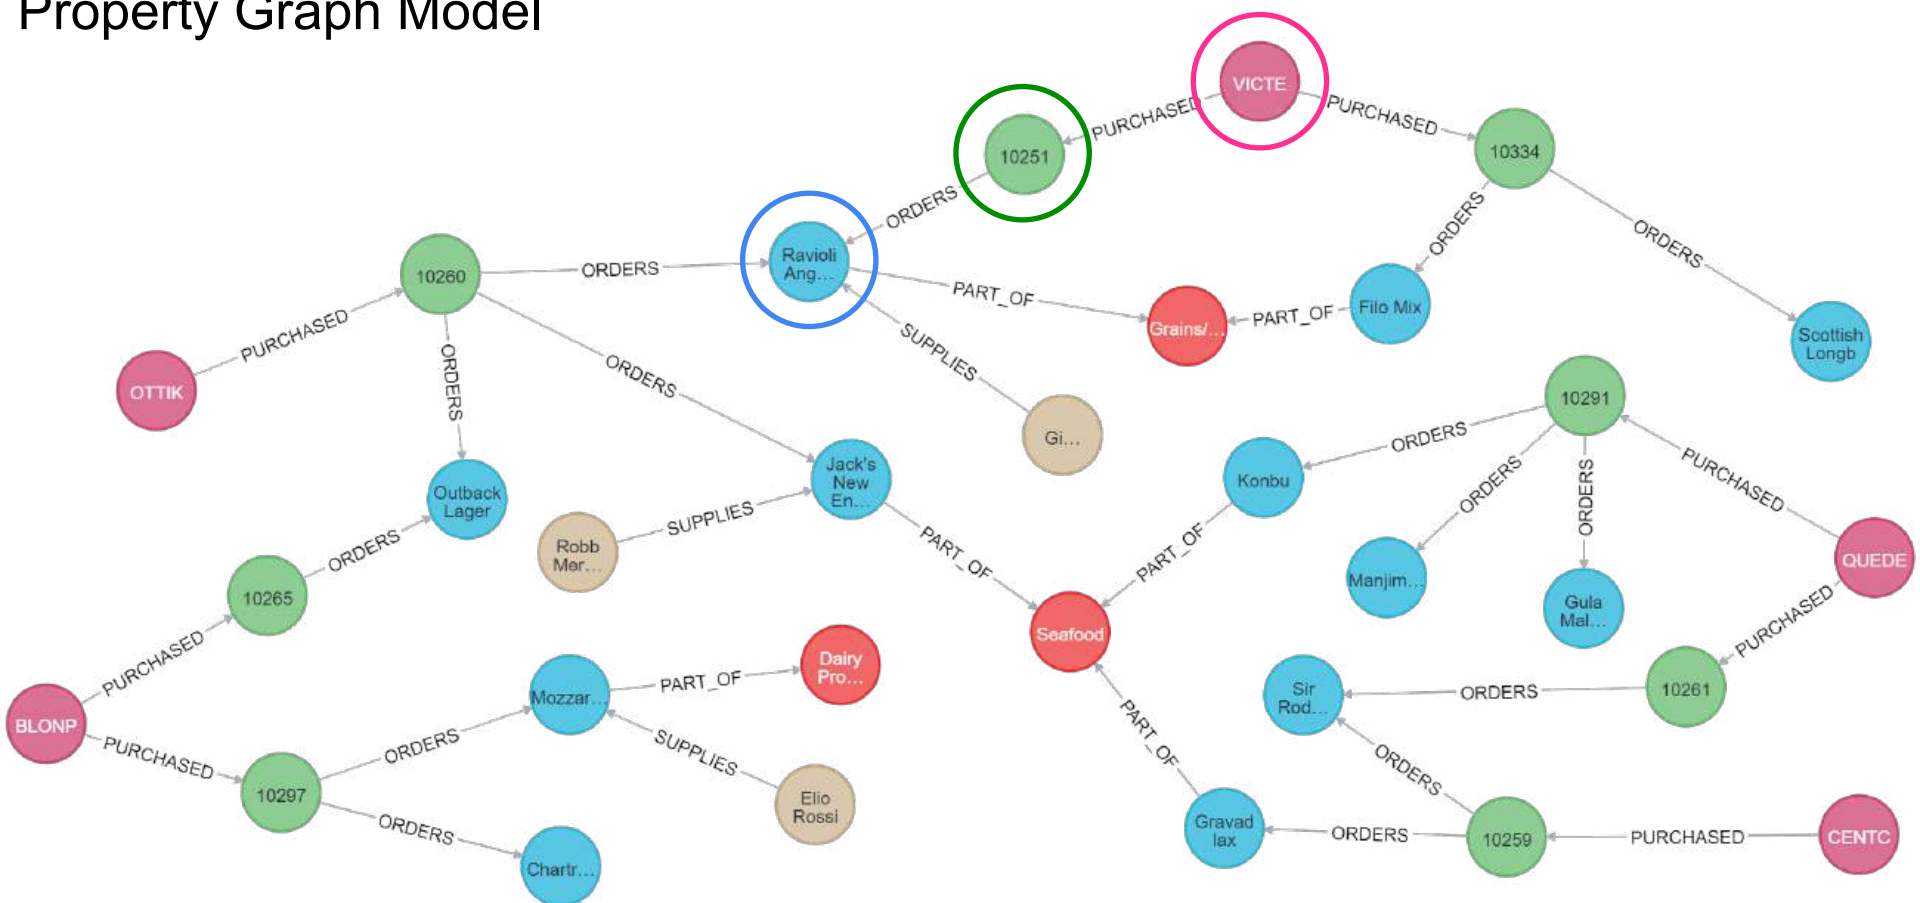
Property Graph Model




Property Graph Model



Property Graph Model



Property Graph Model



Customer Properties

address
city
companyName
contactName
contactTitle
country
customerID

Cus

.....

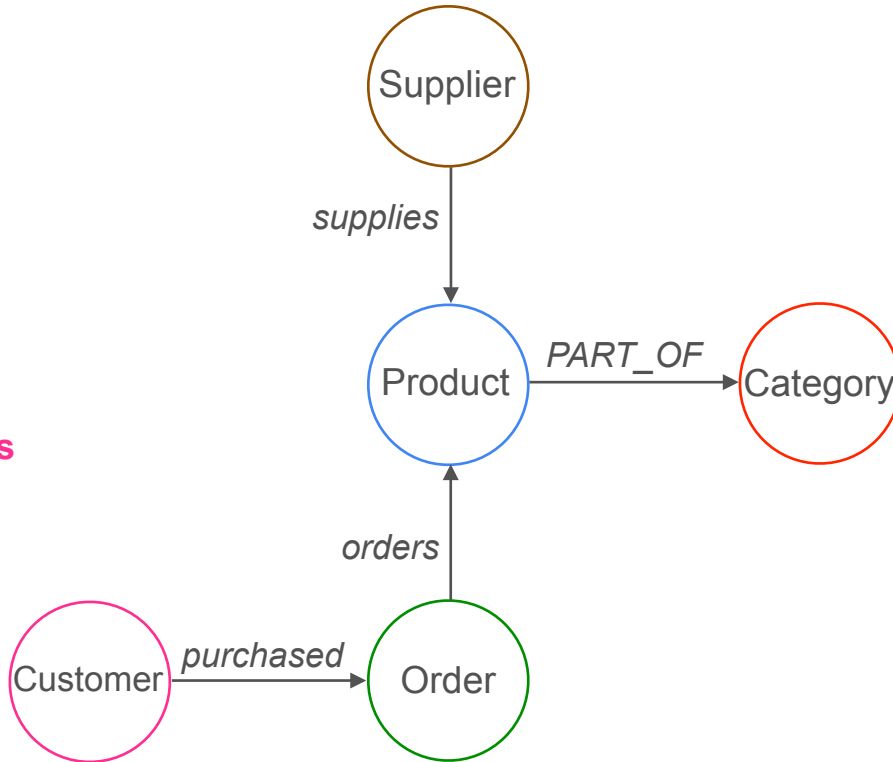
Property Graph Model

Product Properties

productID
productName
unitPrice
unitsInStock
unitsOnOrder

Customer Properties

address
city
companyName
contactName
contactTitle
country
customerID
.....



Supplier Properties

address
city
contactName
fax
region
supplierID
postalCode

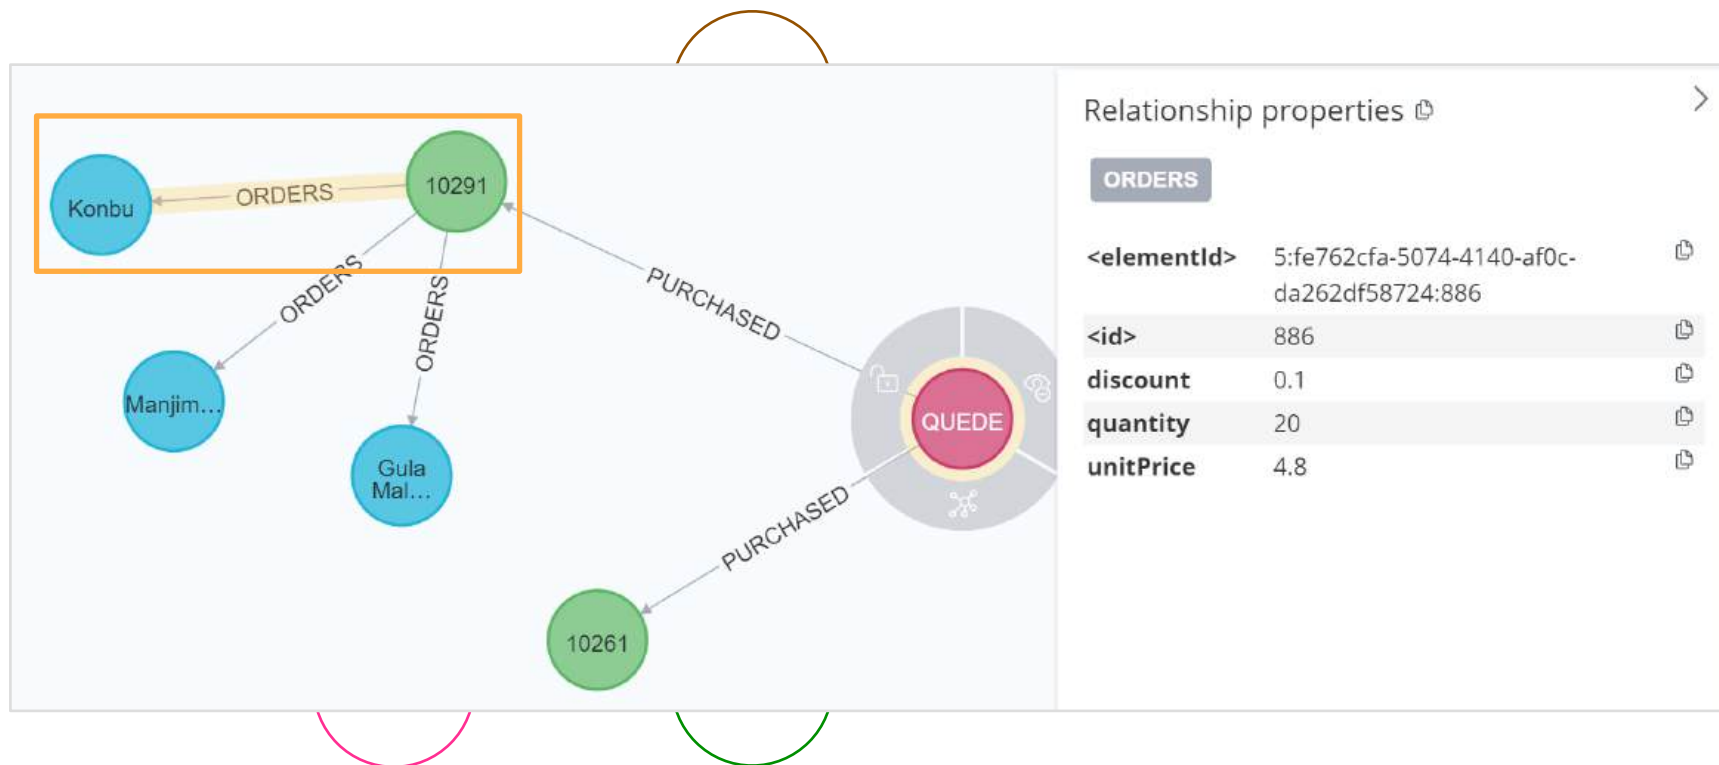
Category Properties

categoryName

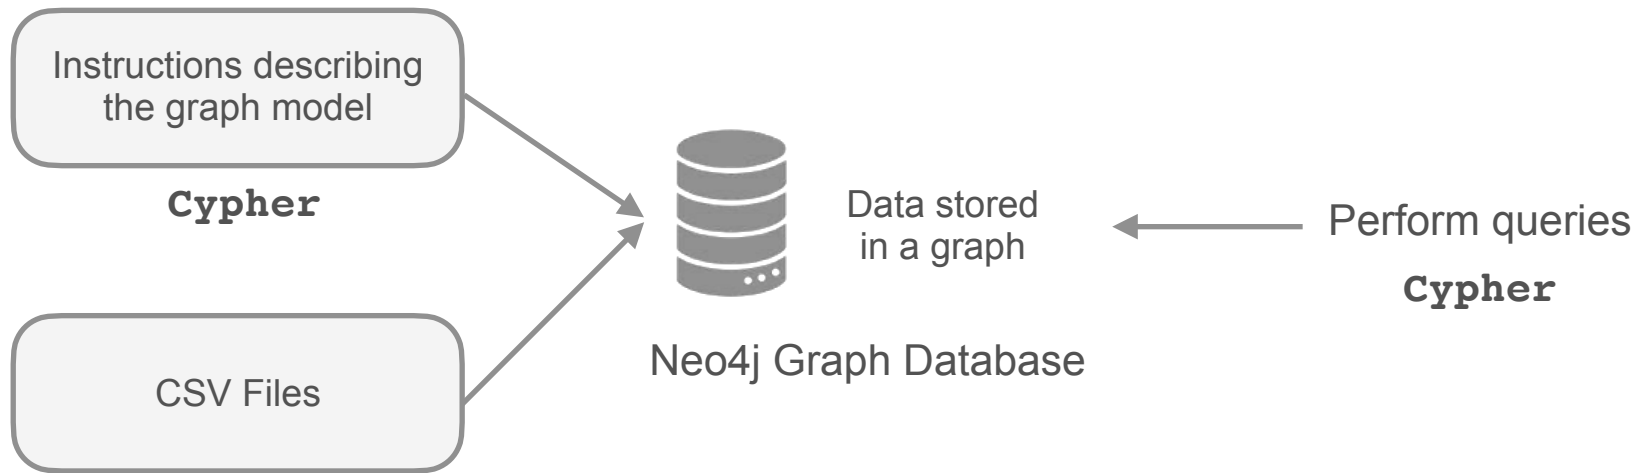
Order Properties

freight
orderDate
orderID
requiredDate
shipAddress
.....

Property Graph Model



Creating a Graph Database



- In the next video, we'll go through some queries examples.
- In the lab, you'll also practice CRUD operations.




DeepLearning.AI

Storage Systems

Neo4j Graph Database & Cypher Query Language (Part 2)

MATCH *pattern* RETURN *result*

node ()




MATCH Statement

Retrieve all nodes



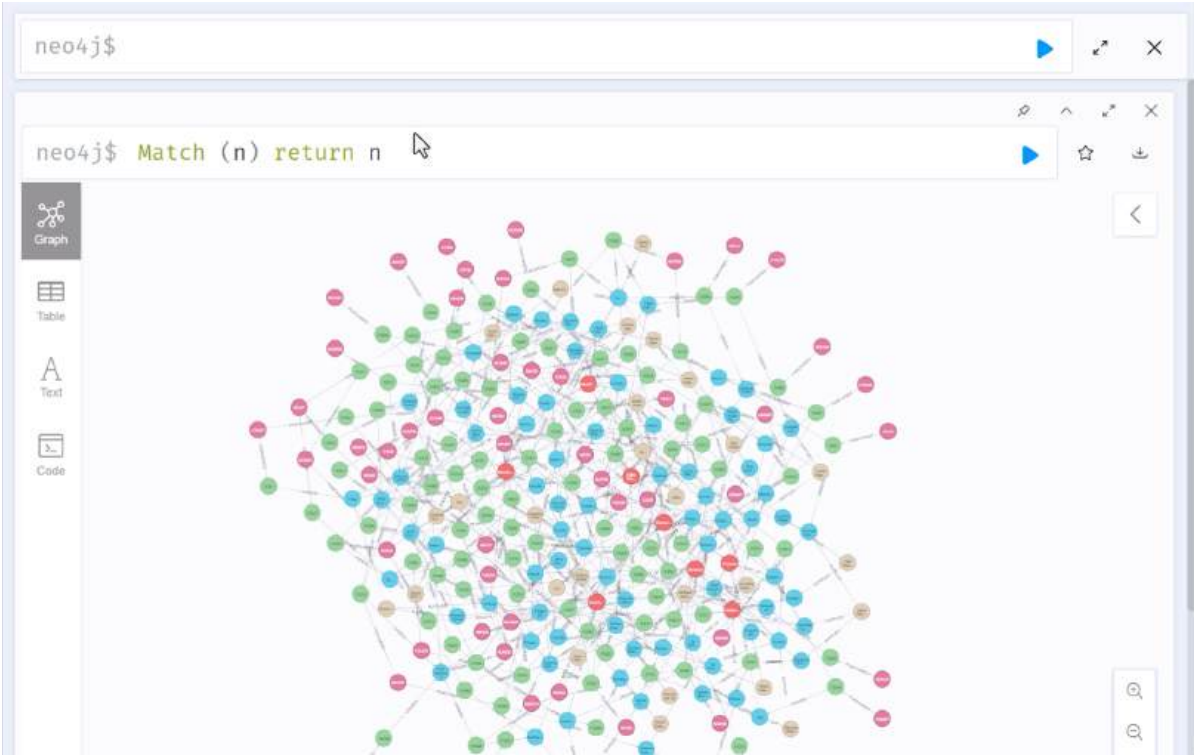
MATCH *pattern* RETURN *result*

node ()



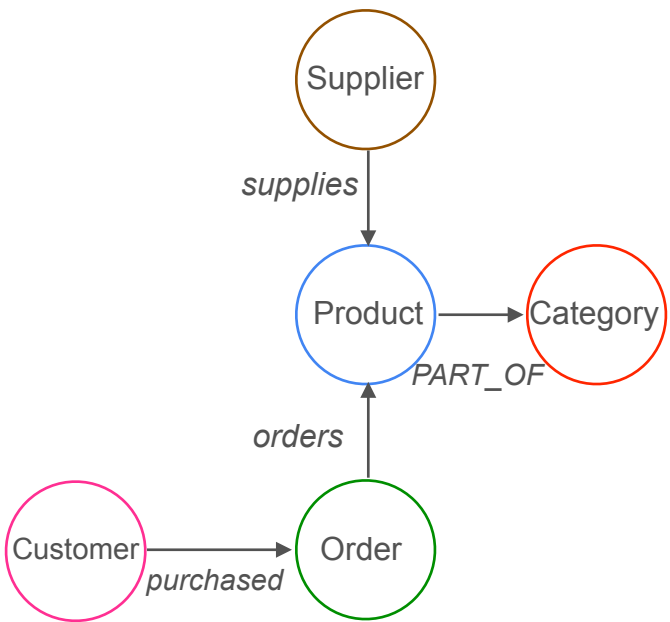
MATCH Statement

Get the total number of nodes



MATCH *pattern* RETURN *result*

node ()



MATCH Statement

Explore the node labels using the `labels` function

```
neo4j$
```

```
neo4j$ Match (n) return count(n)
```


	count(n)
1	265

Started streaming 1 records after 1 ms and completed after 1 ms.

```
neo4j$ Match (n) return n
```

MATCH *pattern* RETURN *result*

node ()



MATCH Statement

Specify the label of the node

neo4j\$


neo4j\$ Match (n) return distinct labels(n)

	labels(n)
1	["Supplier"]
2	["Category"]
3	["Product"]
4	["Order"]
5	["Customer"]

Started streaming 5 records in less than 1 ms and completed after 23 ms.

MATCH *pattern* RETURN *result*

node ()



MATCH Statement

Explore the properties of each order node using the `Properties` function

```
neo4j$
```

```
neo4j$ Match (n:Order) return count(n)
```

	count(n)
1	99


Started streaming 1 records after 2 ms and completed after 10 ms.

```
neo4j$ Match (n) return distinct labels(n)
```

	labels(n)
--	-----------

MATCH *pattern* RETURN *result*

node ()



MATCH Statement

Explore the properties of each order node using the `Properties` function

neo4j\$

neo4j\$ `Match (n:Order) return Properties(n)`

Table

Text

Code

Properties(n)


1

```
{
  "shipCity": "Reims",
  "orderId": "10248",
  "shippedDate": "00:00.0",
  "orderDate": "00:00.0",
  "shipRegion": "NULL",
  "freight": "32.38",
  "shipName": "Vins et alcools Chevalier",
  "shipCountry": "France",
  "shipAddress": "59 rue de l'Abbaye",
  "requiredDate": "00:00.0",
  "shipPostalCode": "51100"
}
```

2

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[r]->(target node)



MATCH Statement

Count all the directed paths

neo4j\$

neo4j\$ Match (n:Order) return Properties(n) limit 1

Table

Properties(n)


1

```
{
  "shipCity": "Reims",
  "orderId": "10248",
  "shippedDate": "00:00.0",
  "orderDate": "00:00.0",
  "shipRegion": "NULL",
  "freight": "32.38",
  "shipName": "Vins et alcools Chevalier",
  "shipCountry": "France",
  "shipAddress": "59 rue de l'Abbaye",
  "requiredDate": "00:00.0",
  "shipPostalCode": "51100"
}
```

Started streaming 1 records after 2 ms and completed after 8 ms.

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[r]->(target node)



MATCH Statement

Return the types of relationships

```
neo4j$
```

```
neo4j$ Match ()-[r]->() return count(r)
```

count(r)	
1	518


Started streaming 1 records in less than 1 ms and completed after 5 ms.

```
neo4j$ Match (n:Order) return Properties(n) limit 1
```

Properties(n)	
---------------	--

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[r]->(target node)



MATCH Statement

Specify the type of the relationship

```
neo4j$
```

```
neo4j$ Match ()-[r]->() return distinct type(r)
```


	type(r)
1	"SUPPLIES"
2	"PART_OF"
3	"PURCHASED"
4	"ORDERS"

Started streaming 4 records after 1 ms and completed after 21 ms.

```
neo4j$ Match (n:Order) return Properties(n) limit 1
```

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[r]->(target node)



MATCH Statement

Return the properties of a relationship

```
neo4j$ Match ()-[r:ORDERS]->()
```

```
neo4j$ Match ()-[r]->() return distinct type(r)
```

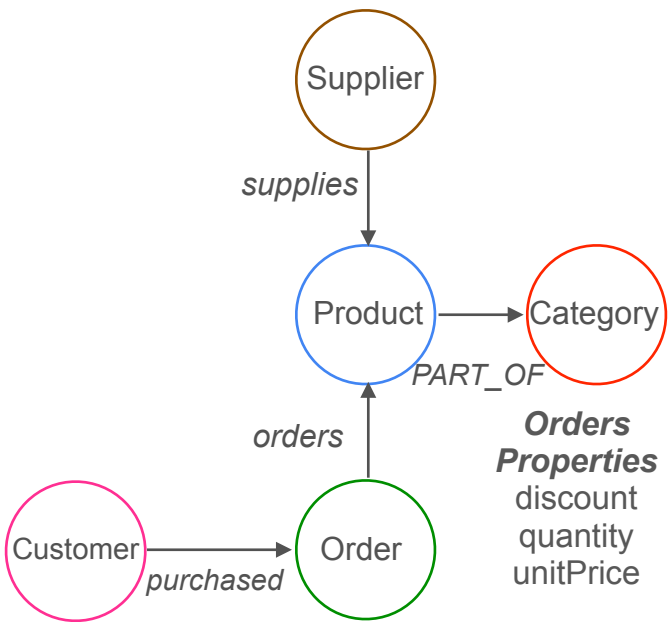
	type(r)
1	"SUPPLIES"
2	"PART_OF"
3	"PURCHASED"
4	"ORDERS"

Started streaming 4 records after 1 ms and completed after 21 ms.

```
neo4j$ Match (n:Order) return Properties(n) limit 1
```

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Return the properties of a relationship

```
neo4j$
```

```
neo4j$ Match ()-[r:ORDERS]->() return AVG(r.quantity*r.unitPrice)
```

Table	AVG(r.quantity*r.unitPrice)
1	502.4645283018867


Started streaming 1 records after 1 ms and completed after 25 ms.

```
neo4j$ Match ()-[r]->() return distinct type(r)
```

Table	type(r)
-------	---------

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Get the average price for all orders grouped by product category

```
neo4j$
```

```
neo4j$ Match ()-[r:ORDERS]->() return AVG(r.quantity*r.unitPrice) as a...
```

	average_price
1	502.4645283018867


Started streaming 1 records in less than 1 ms and completed after 25 ms.

```
neo4j$ Match ()-[r]->() return distinct type(r)
```

type(r)

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Get the average price for all orders grouped by product category

```
neo4j$ Match ()-[r:ORDERS]->()-[part:PART_OF]->(c:Category) return  
AVG(r.quantity*r.unitPrice) as average_price
```

```
neo4j$ Match ()-[r:ORDERS]->() return AVG(r.quantity*r.unitPrice) as a...
```

	average_price
1	502.4645283018867


Started streaming 1 records in less than 1 ms and completed after 25 ms.

```
neo4j$ Match ()-[r]->() return distinct type(r)
```

type(r)

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”


neo4j\$

neo4j\$ Match ()-[r:ORDERS]->()-[part:PART_OF]->(c:Category) return c.c...

	c.categoryName	average_price
2	"Confections"	531.36666666666664
4	"Dairy Products"	504.31836734693877
5	"Grains/Cereals"	278.67
6	"Meat/Poultry"	766.9142857142859
7	"Produce"	645.6571428571427
8	"Seafood"	434.8282051282051

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”

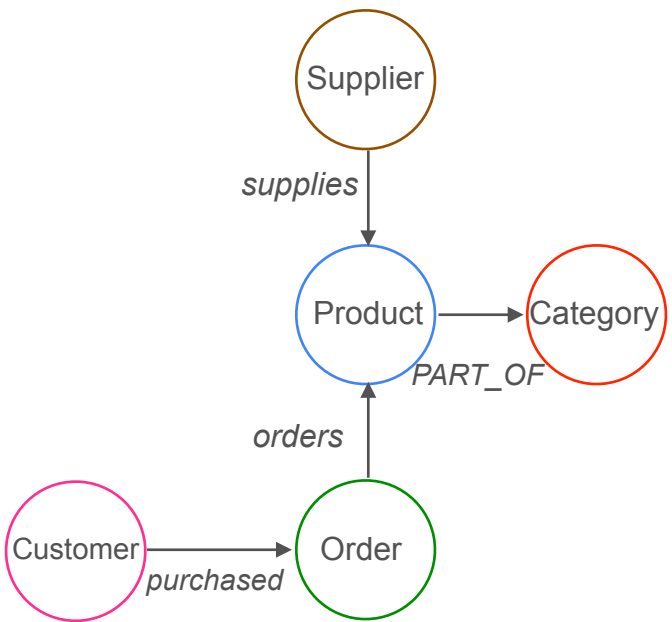
```
neo4j$ Match (p:Product)-[:PART_OF]->(c:Category)
```

```
neo4j$ Match ()-[r:ORDERS]->()-[part:PART_OF]->(c:Category) return c.c...
```

	c.categoryName	average_price
2	"Confections"	531.36666666666664
4	"Dairy Products"	504.31836734693877
5	"Grains/Cereals"	278.67
6	"Meat/Poultry"	766.9142857142859
7	"Produce"	645.6571428571427
8	"Seafood"	434.8282051282051

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”

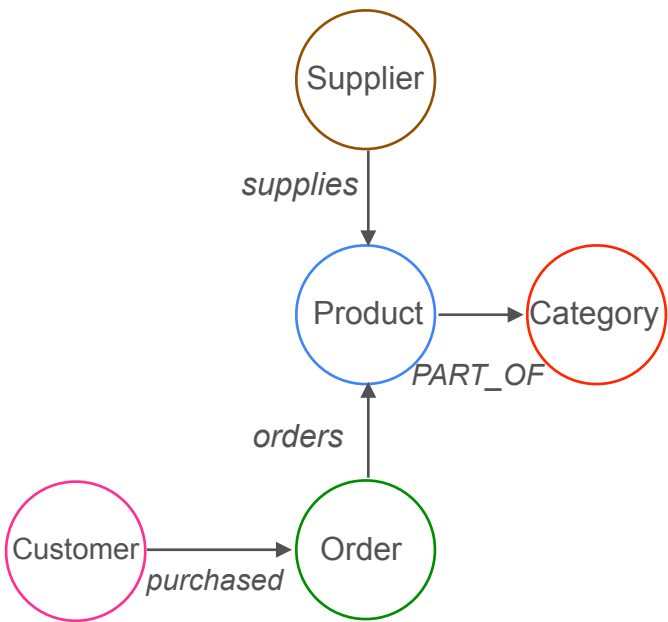
```
1 Match (p:Product)-[:PART_OF]→(c:Category)
2 where c.categoryName="Meat/Poultry"
```

```
neo4j$ Match ()-[r:ORDERS]→()-[part:PART_OF]→(c:Category) return c.c...
```

	c.categoryName	average_price
3	"Confections"	531.36666666666664
4	"Dairy Products"	504.31836734693877
5	"Grains/Cereals"	278.67
6	"Meat/Poultry"	766.9142857142859
7	"Produce"	645.6571428571427
8	"Seafood"	434.8282051282051

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”

neo4j\$


```
1 Match (p:Product)-[:PART_OF]-(c:Category)
2 where c.categoryName="Meat/Poultry"
3 return p.productName, p.unitPrice
```

Table

	p.productName	p.unitPrice
1	"Perth Pasties"	32.8
2	"Tourtière"	7.45
3	"Alice Mutton"	39.0
4	"Pâté chinois"	24.0
5	"Thüringer Rostbratwurst"	123.79
6	"Mishi Kobe Niku"	97.0

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Retrieve the product name of all products ordered by the customer “QUEDE”

neo4j\$


neo4j\$ Match (p:Product)-[:PART_OF]-(c:Category {categoryName:"Meat/P...})

Table

	p.productName	p.unitPrice
1	"Perth Pasties"	32.8
2	"Tourtière"	7.45
3	"Alice Mutton"	39.0
4	"Pâté chinois"	24.0
5	"Thüringer Rostbratwurst"	123.79
6	"Mishi Kobe Niku"	97.0

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Get the ID of other customers who ordered the same products as “QUEDE”

```
neo4j$ Match (c1:Customer {customerID:"QUEDE"}) -[:PURCHASED]->()-[:ORDERS]->(p:Product)
```


```
neo4j$ Match (c1:Customer {customerID:"QUEDE"}) -[:PURCHASED]->()-[:OR...
```

	p.productName
1	"Steeleye Stout"
2	"Sir Rodney's Scones"
3	"Gula Malacca"
4	"Konbu"
5	"Manjimup Dried Apples"

Started streaming 5 records after 1 ms and completed after 1 ms.

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Retrieve the orders that contain at most two products


neo4j\$

neo4j\$ Match (c1:Customer {customerID:"QUEDE"}) -[:PURCHASED]->()-[:OR...

	c2.customerID
1	"QUICK"
2	"SAVEA"
3	"ROMEY"
4	"ISLAT"
5	"WARTH"
6	"CENTC"

MATCH *pattern* RETURN *result*

node	()
relationship	[]
path	(source node)-[]->(target node)



MATCH Statement

Retrieve the orders that contain at most two products

neo4j\$

```
1 Match (o:Order)-[:ORDERS]-(p:Product)
2 return o.orderID as ID, count(p) as countProd
```

ID	countProd
"10294"	5
"10317"	1
"10285"	3
"10342"	4
"10255"	4
"10327"	4



DeepLearning.AI

Storage Systems

Summary

Raw Storage Ingredients

Persistent Storage Medium

Magnetic disk




Solid-state storage



Volatile Memory


RAM




CPU cache

Cloud Storage Options

File Storage



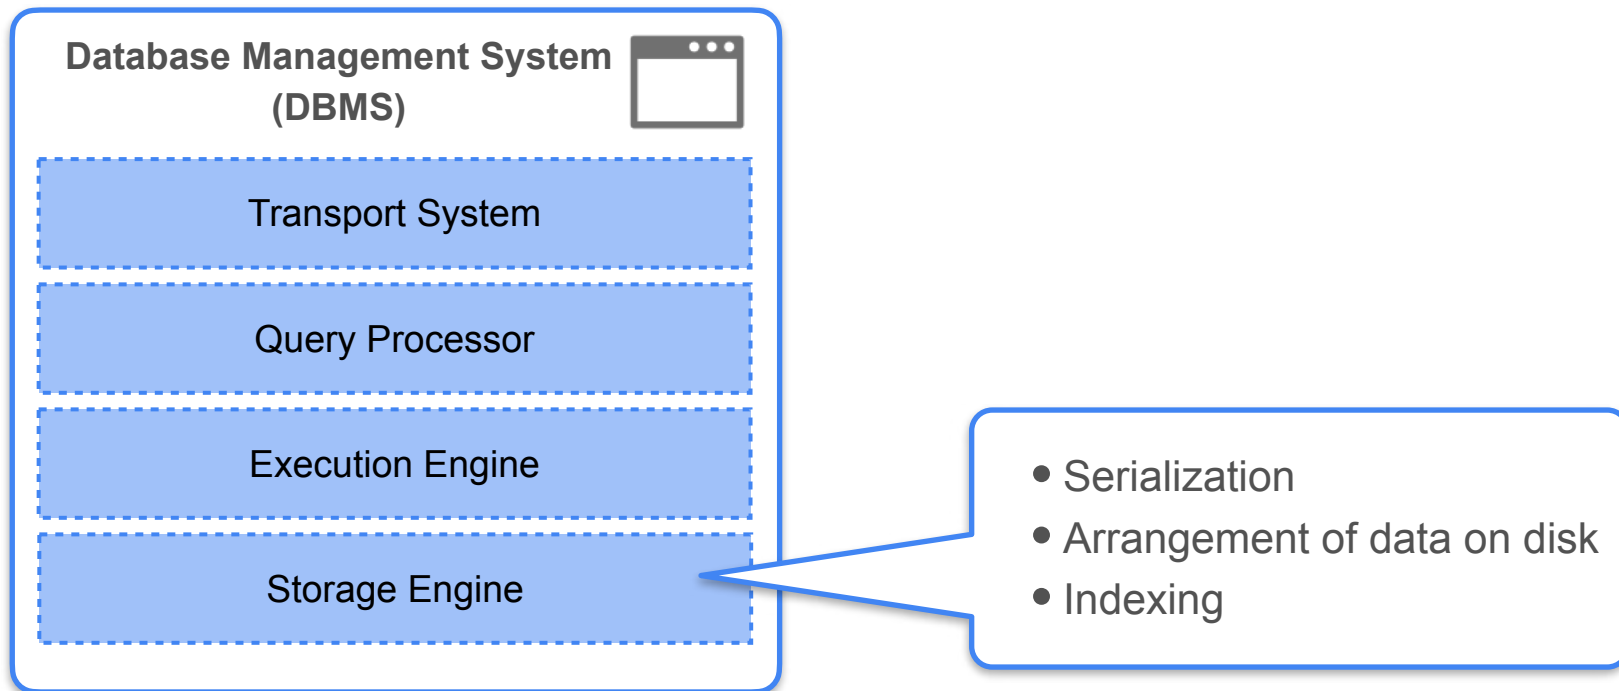
Block Storage



Object Storage



Storage in Databases



Row and Columnar Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	458650	10	67t
2	23	902348	14	56t
3	45	1255893	12	87q
4	50	456829	13	98q

Row-oriented storage


bytes representing the 1st row	bytes representing the 2nd row	...	bytes representing the last row
--------------------------------	--------------------------------	-----	---------------------------------

Column-oriented storage

bytes representing the 1st column	bytes representing the 2nd column	...	bytes representing the last column
-----------------------------------	-----------------------------------	-----	------------------------------------

Databases

Graph Databases



Cypher

Vector Databases

